

# CONTAINERIZATION TRAINING

# HOW WE TEACH

// We believe in **learning by doing**

// The training is **lab driven**

// Work together!

// Ask questions at any time

# SESSION LOGISTICS

// 2 days duration

// Mostly exercises

// Regular breaks

# ASSUMED KNOWLEDGE AND REQUIREMENTS

// Familiarity with Bash or Powershell

// Bash Cheat Sheet: <http://bit.ly/2mTQr8l>

# YOUR LAB ENVIRONMENT

// You have been given an instance for use in exercises

// Ask instructor for credentials if you don't have them already

# COURSE LEARNING OBJECTIVES

By the end of this course, learners will be able to

// Assess the advantages of a containerized software development & deployment

// Use Docker engine features necessary for running containerized applications

# INTRODUCING DOCKER

# WHAT WE WANT

Ideal software should

// be modular and flexible (devs)

// be easy to migrate (devops)

// be easy to scale, monitor and lifecycle (ops)

// mitigate vulnerabilities (security)

// run cheap (business)



# WITHOUT CONTAINERIZATION

DB + API + Frontend

DB Config + API Config + Frontend Config

DB Dependencies + API Dependencies + Frontend Dependencies

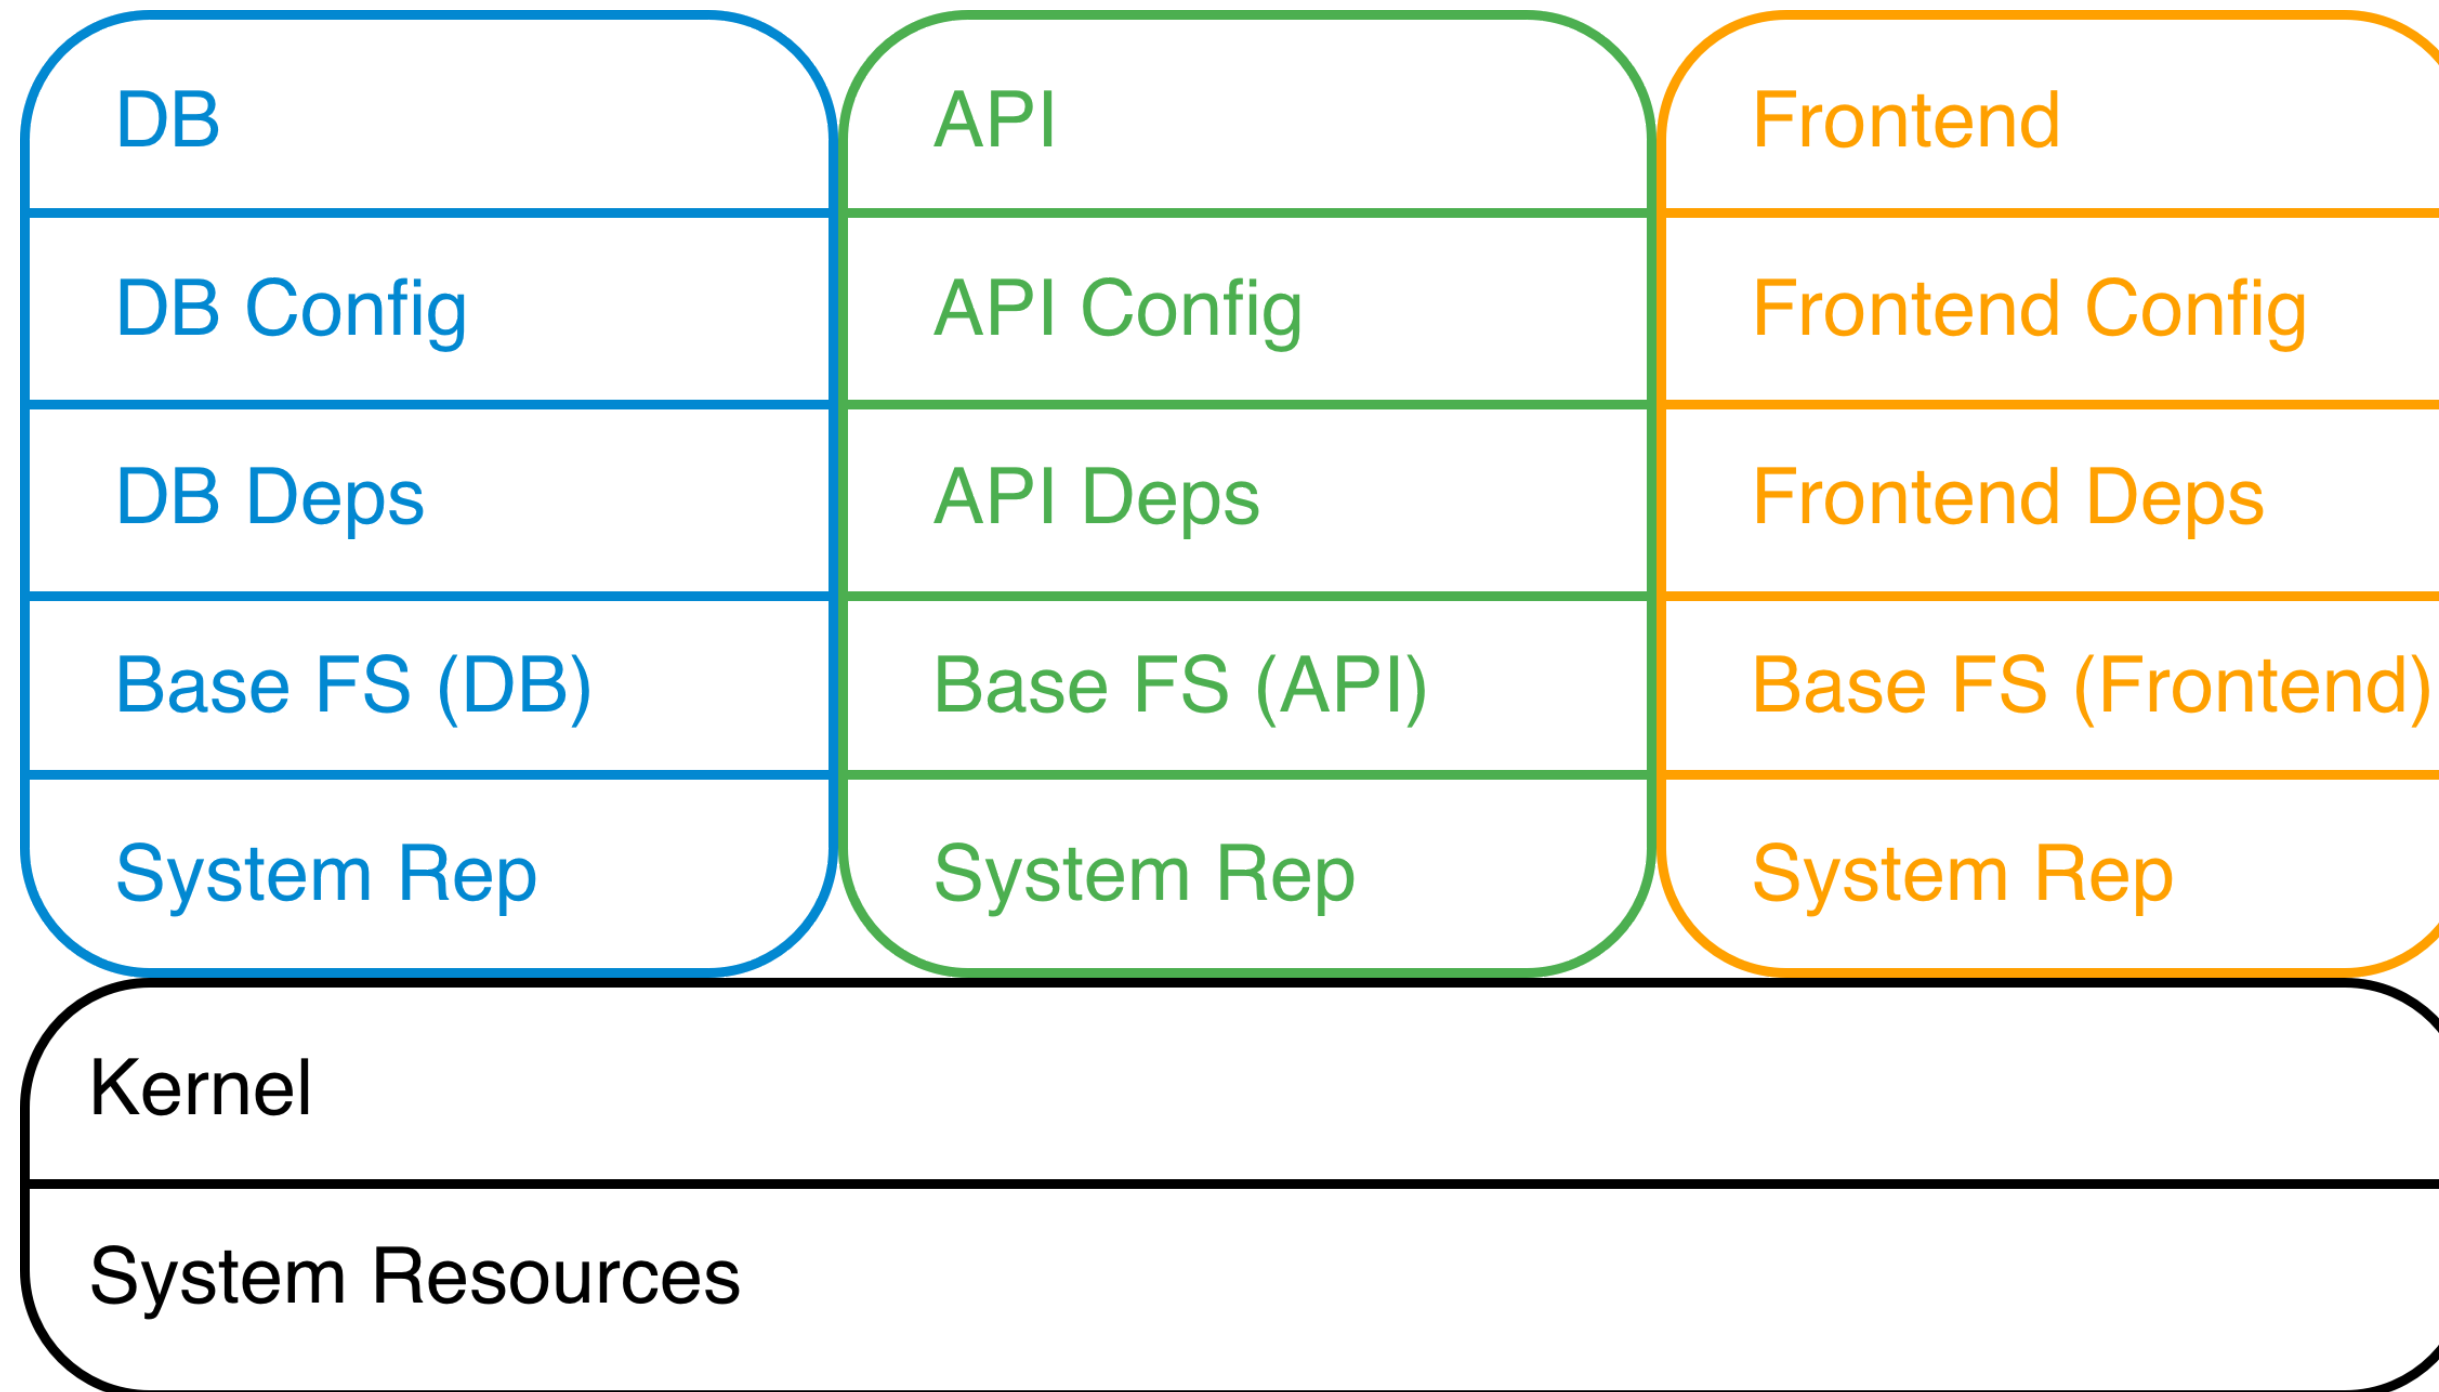
Base Filesystem

System Representation (FS, ports, process tree, users...)


Kernel

System Resources

# WITH CONTAINERIZATION



# RAPID DEVELOPMENT




Containers can be removed and replaced with a minimum of impact on their neighbors, increasing developer choice and speed.

# SMOOTH MIGRATION



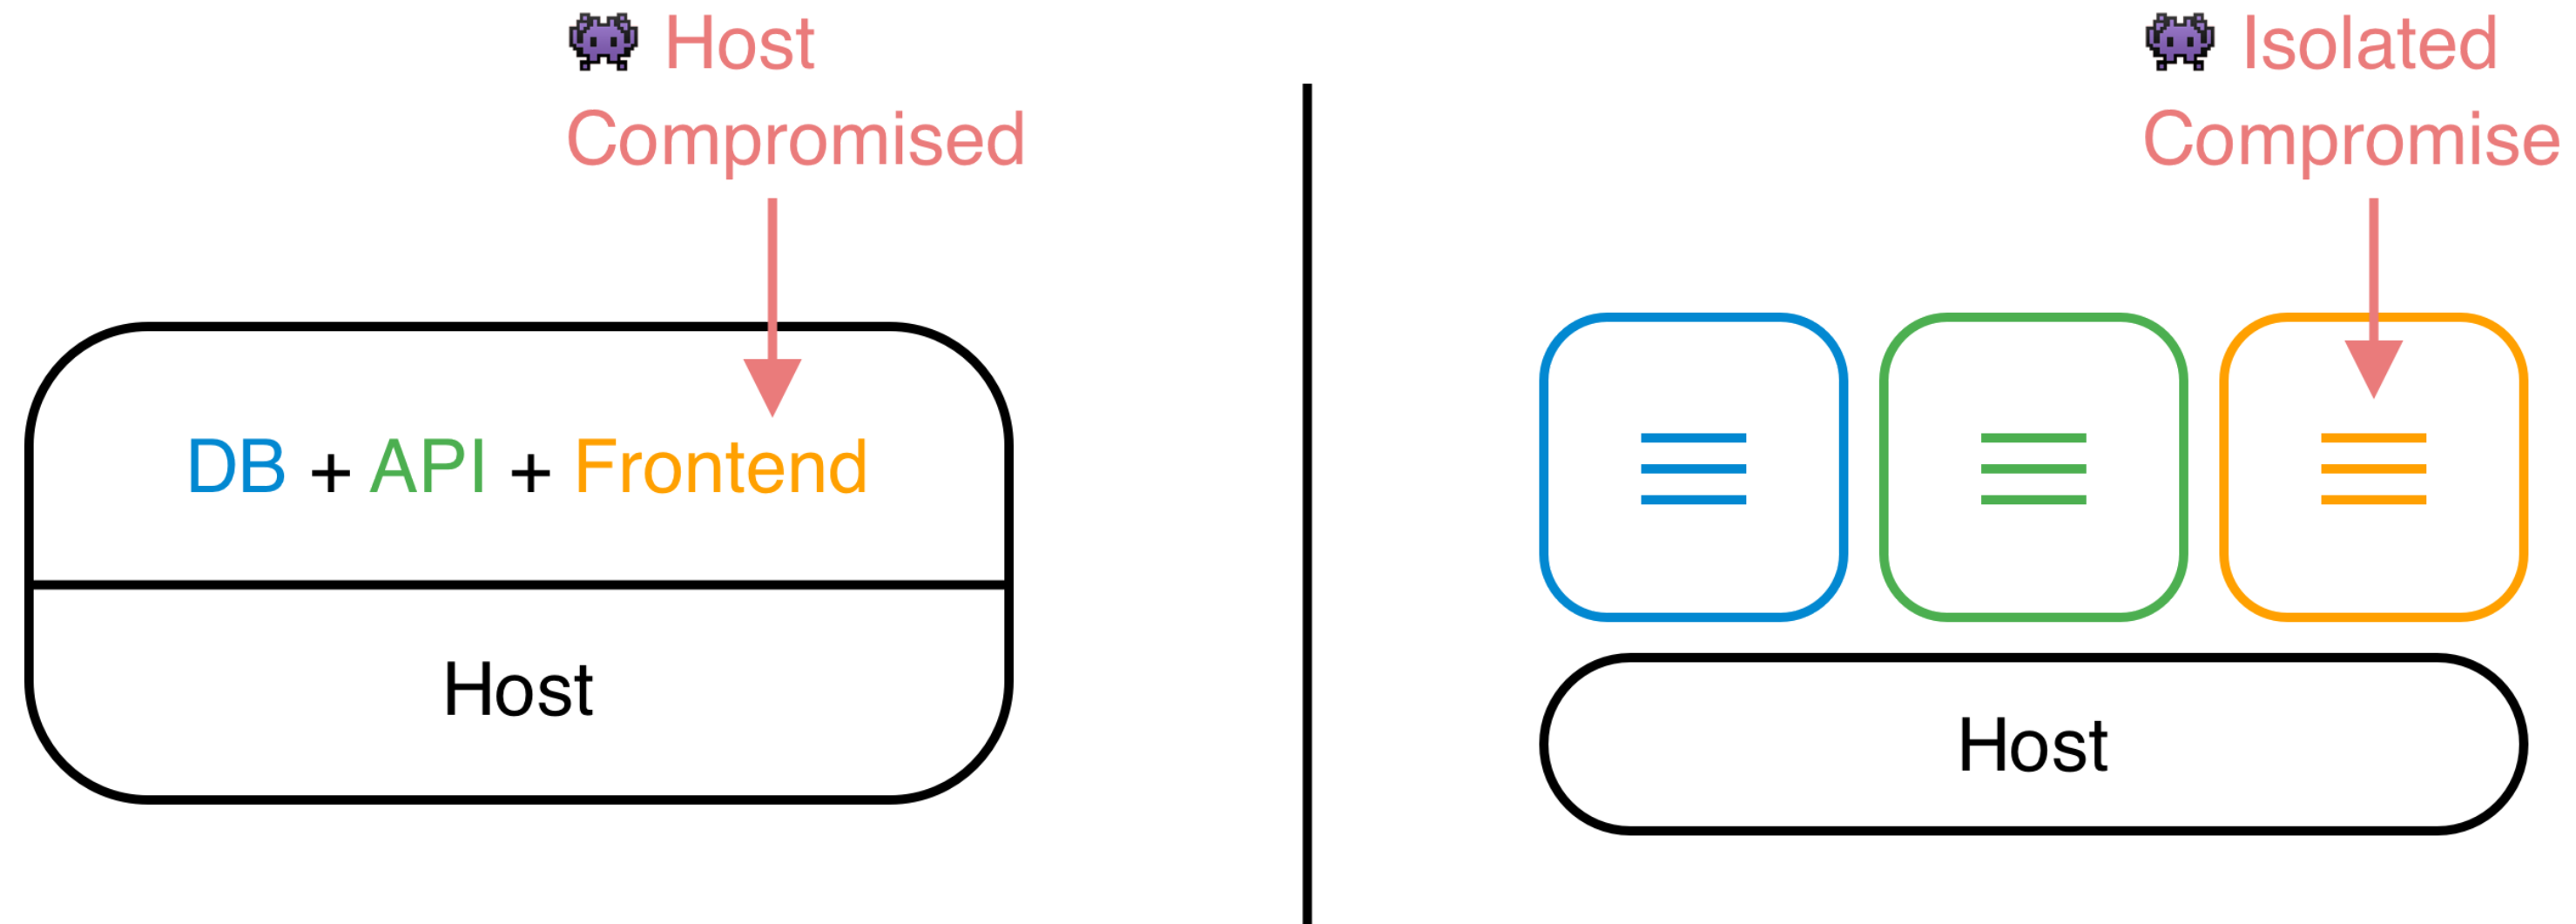
Containers carry their environment and dependencies with them, simplifying and minimizing requirements on the hosts that run them.

# SIMPLE SCALE & MAINTENANCE




Weak coupling between containers minimizes side effects when scaling and simplifies monitoring.

# SECURE BY DEFAULT




Containers have private system resources, so a compromise in one does not affect the rest.

# APPLICATION DENSITY



Virtual Machines



Containers

Containers save datacenter costs by running many more application instances than virtual machines can on the same physical hosts.

# CONTAINERIZATION BASICS



# DISCUSSION: RUNNING CONTAINERS

What assurances would you need to run a process on an arbitrary host? Consider

// Hostile environments

// Required resources

# LEARNING OBJECTIVES

By the end of this module, learners will be able to

// Describe what a container is in terms of processes and isolation tools

// Use the key commands for interacting with Docker containers

# CONTAINERS ARE PROCESSES

**Containers** are processes sandboxed by

// Kernel namespaces

// Control Groups

// Root privilege management & syscall restrictions (Linux)

// VM isolation (Windows)

# LINUX KERNEL NAMESPACES

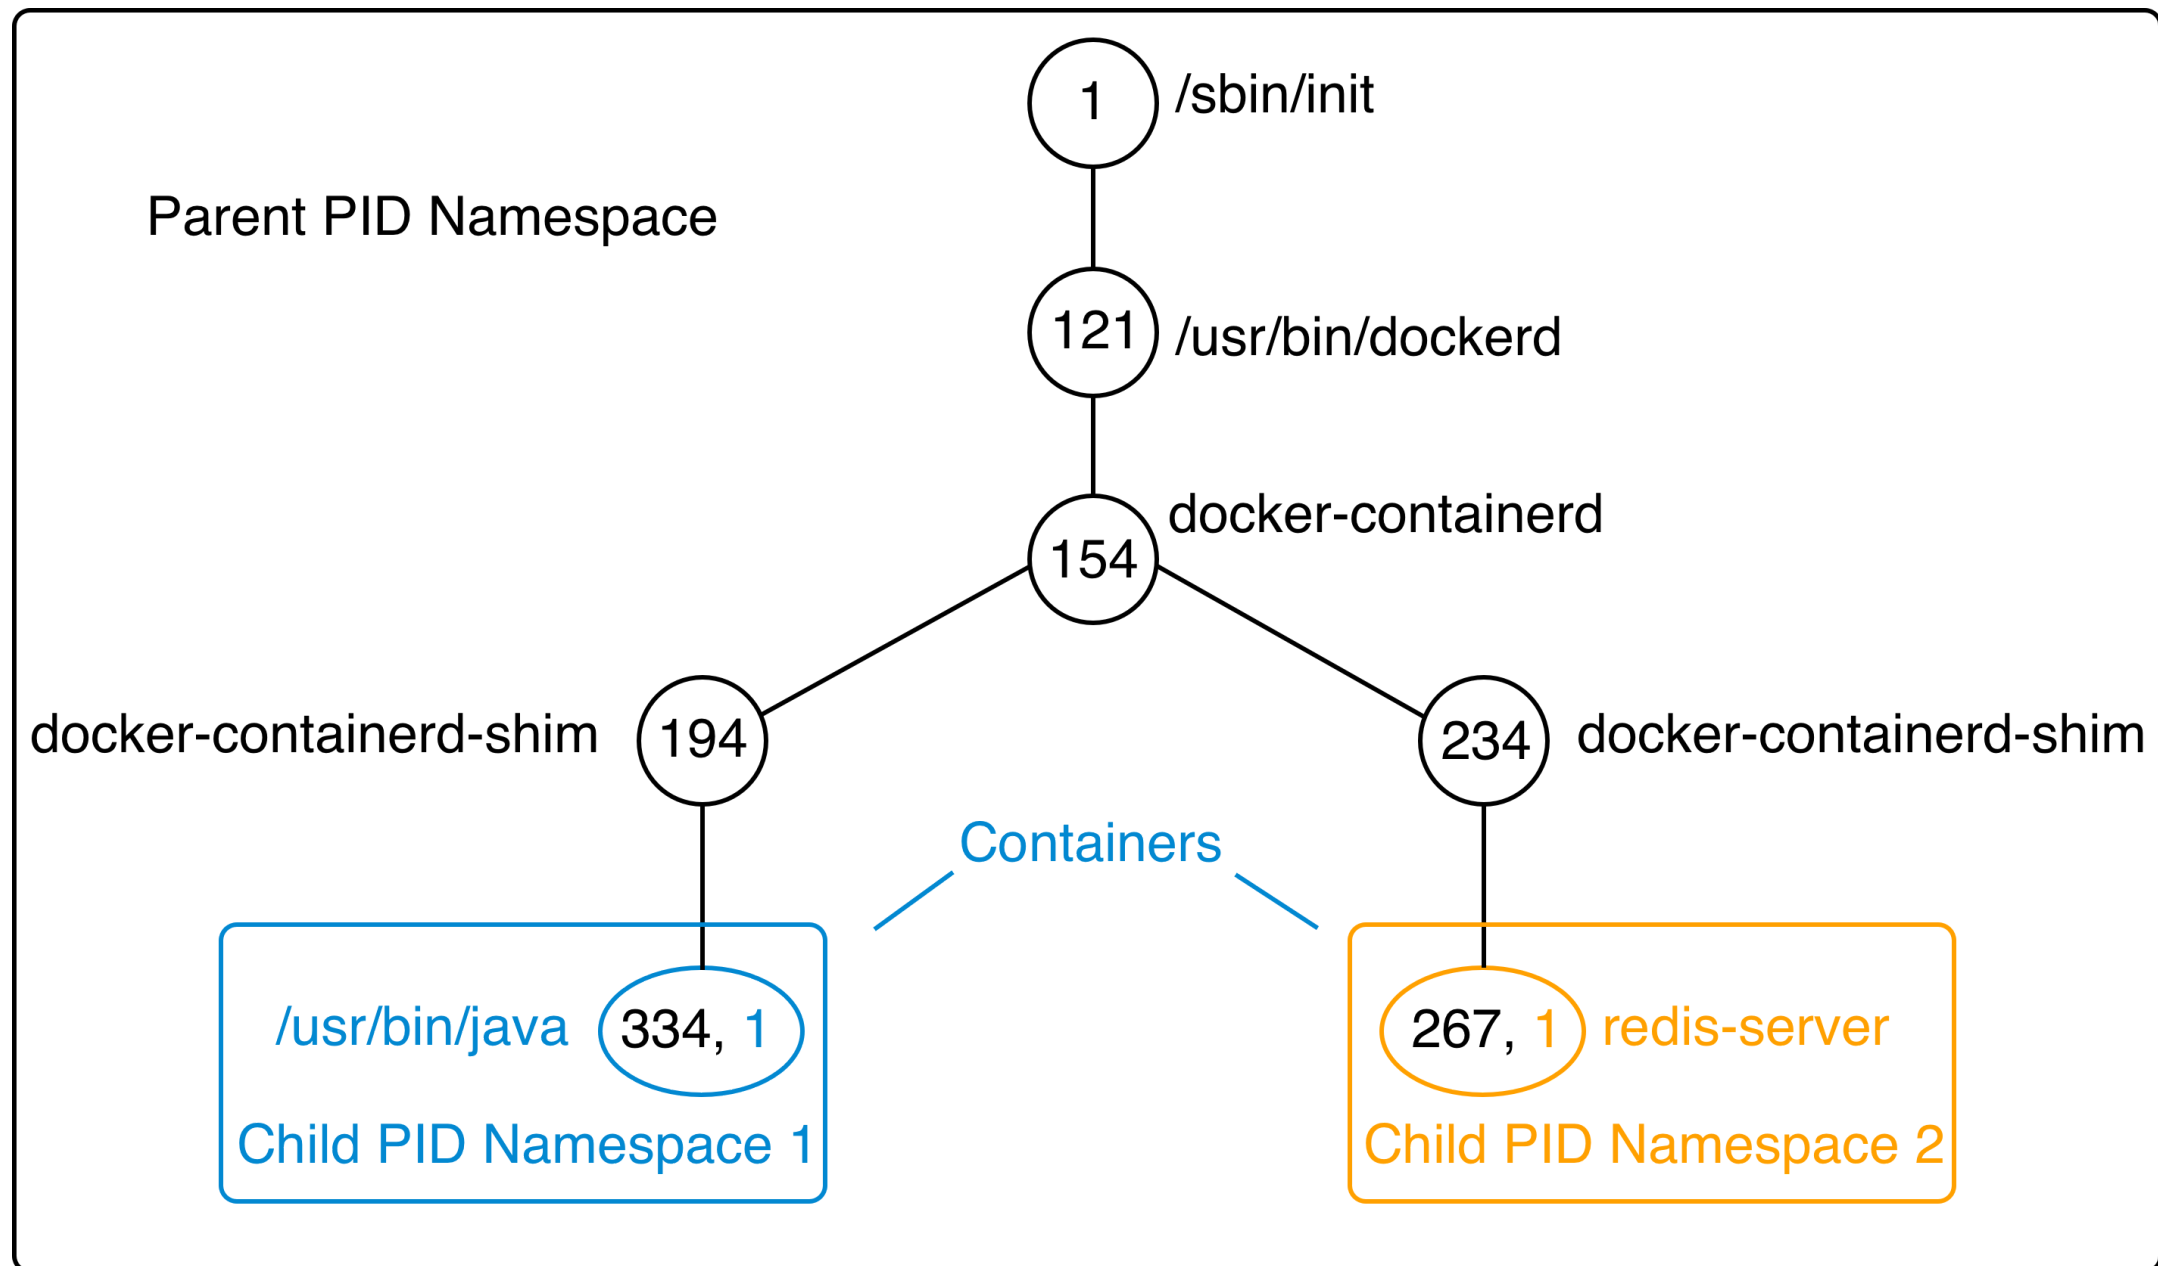
## // **DEFAULT**

- Process IDs
- Network stacks
- Inter-process communications
- Mount points
- Hostnames

## // **OPTIONAL**

- User IDs

# LINUX PID KERNEL NAMESPACE



# OPTIONAL LINUX ISOLATION FEATURES

// Control groups: limit memory & CPU

// Root privilege management: whitelist root powers

// System call management: whitelist available system calls

// Linux Security Modules: mandatory filesystem access control



# INSTRUCTOR DEMO: PROCESS ISOLATION

See the demo

// Process Isolation

In the Exercises book.



## EXERCISE: CONTAINER BASICS

Work through

// Running and Inspecting a Container

// Interactive Containers

// Detached Containers and Logging

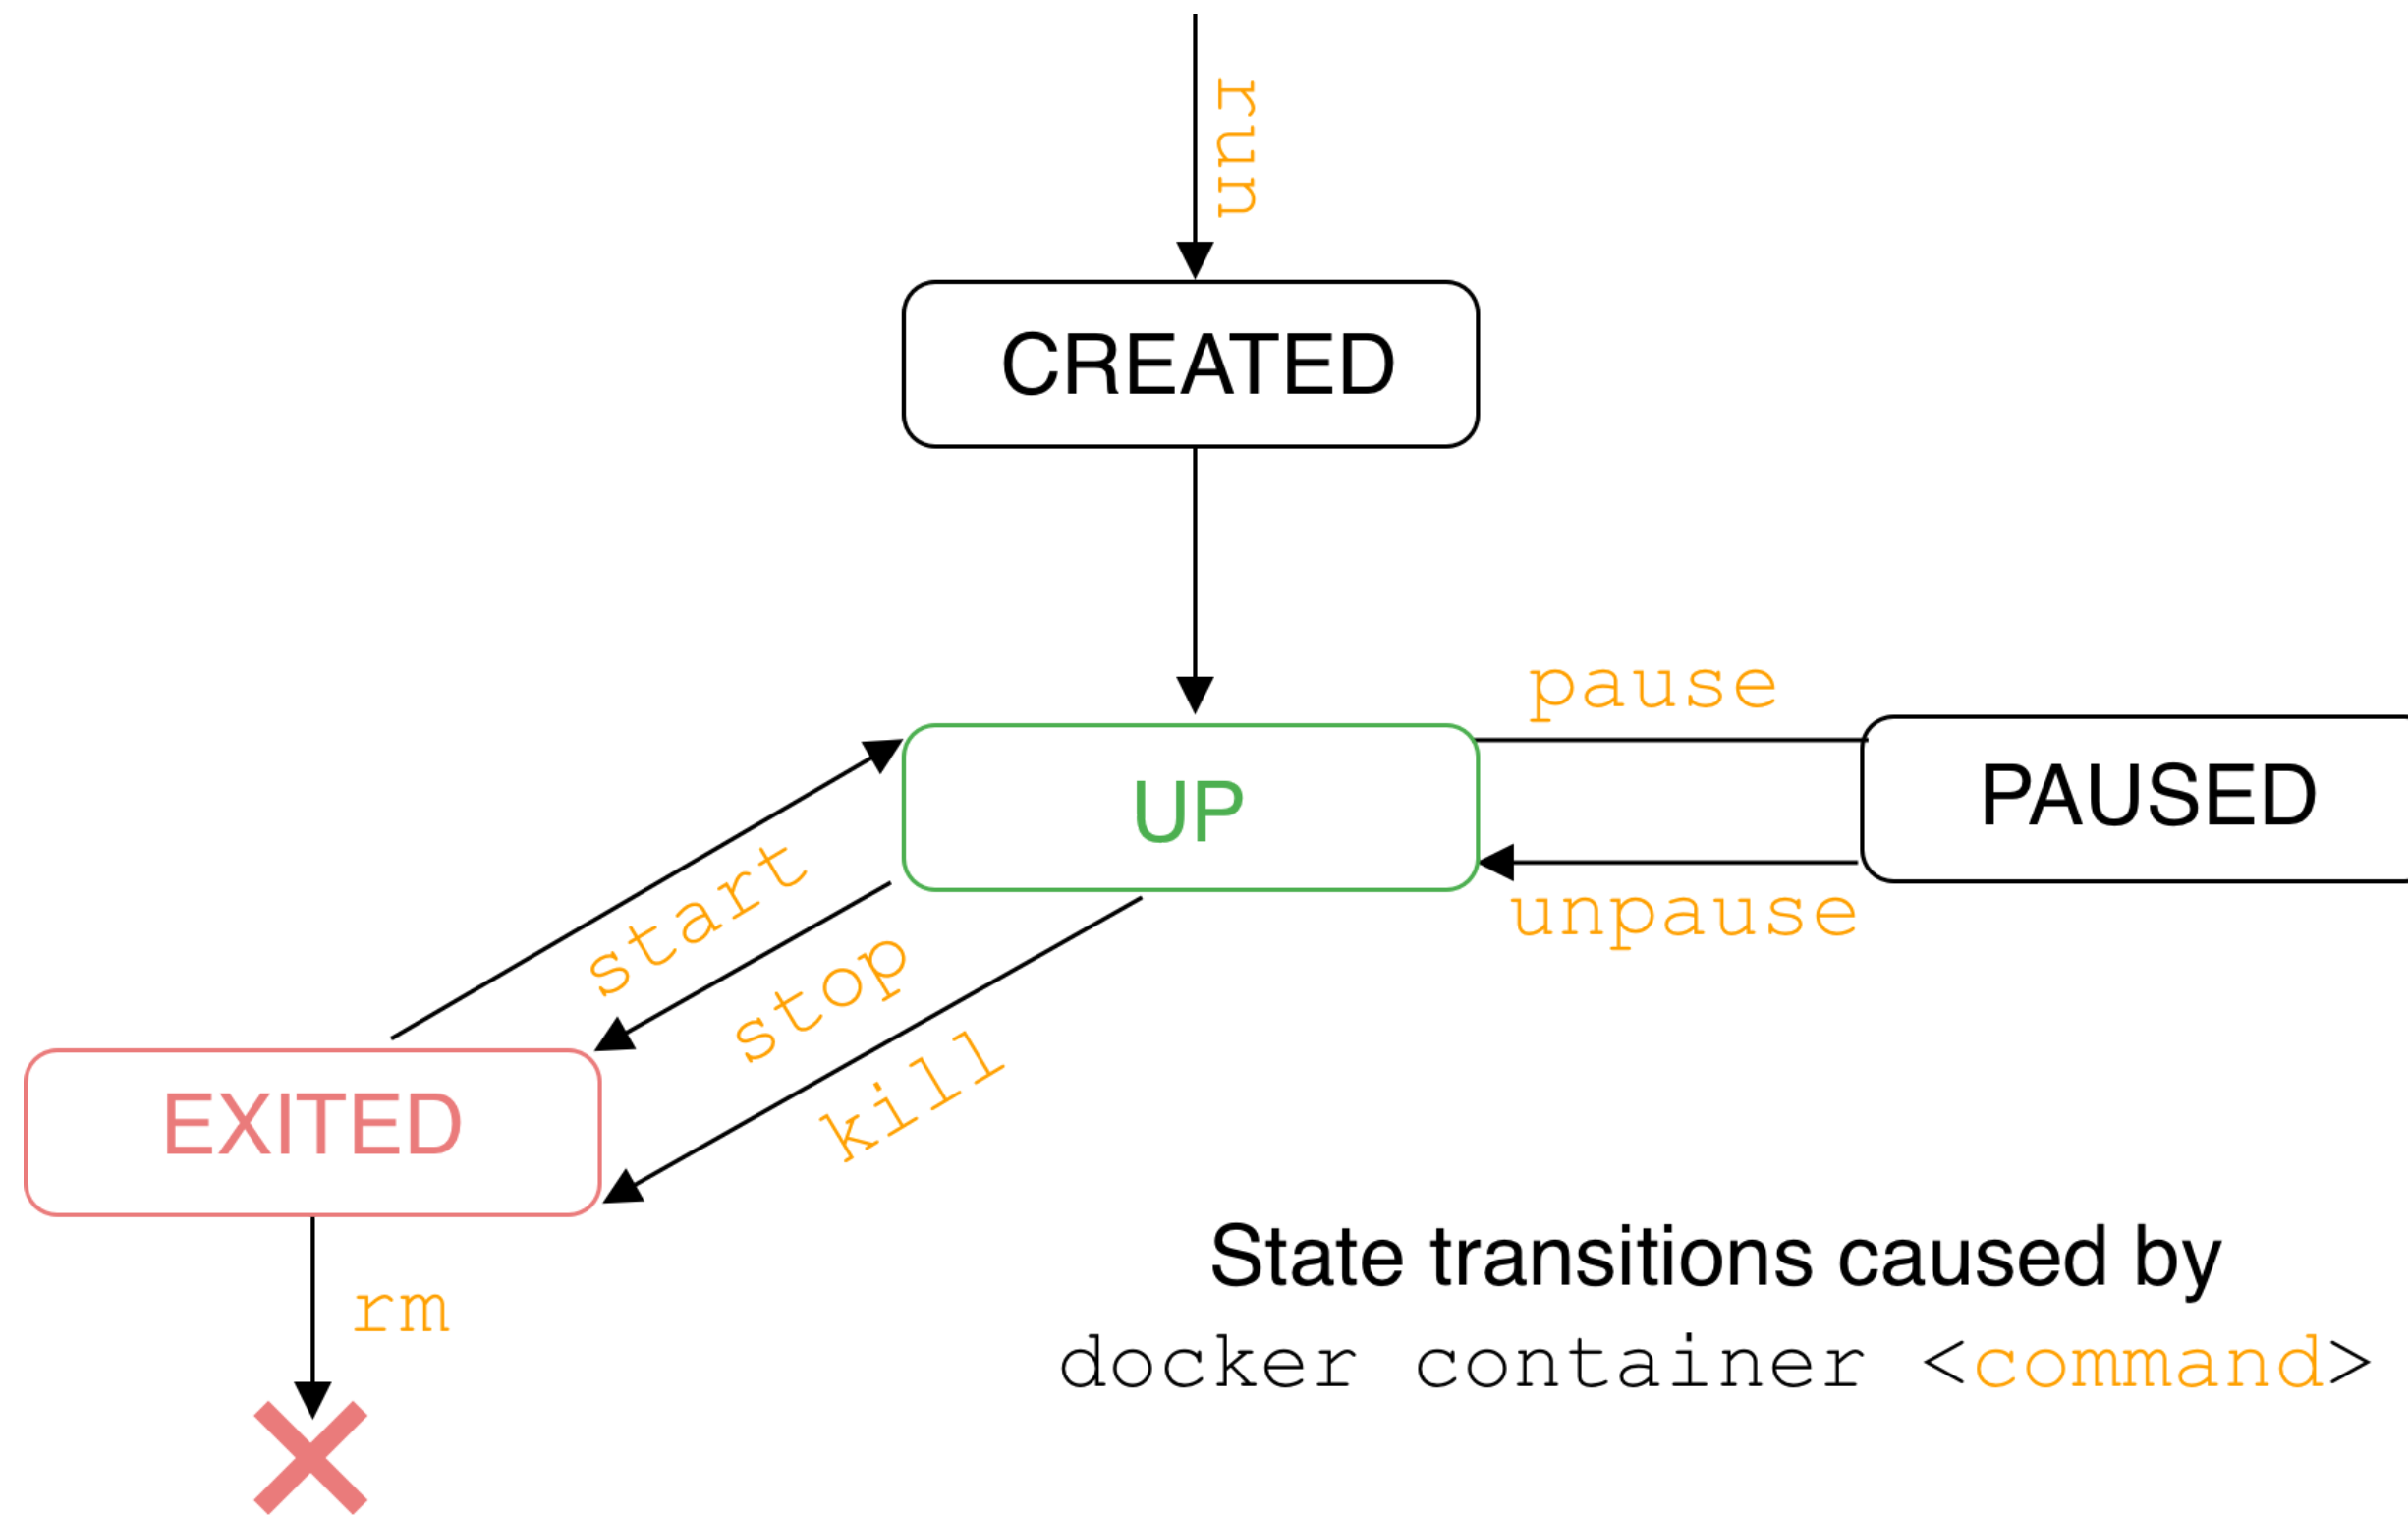
// Starting, Stopping, Inspecting and Deleting Containers

In the Exercises book.

21M 4S



# CONTAINER LIFECYCLE



# CONTAINER LOGS

// STDOUT and STDERR for a containerized process

// `docker container logs <container name>`

# CONTAINER BASICS TAKEAWAYS

// Single process constrained by kernel namespaces, control groups and other technologies

// Private & ephemeral filesystem and data

# FURTHER READING

// List of container commands: <http://dockr.ly/2iLBV2l>

// Getting started with containers: <http://dockr.ly/2gmxKWB>

// Start containers automatically: <http://dockr.ly/2xB8sMI>

// Limit a container's resources: <http://dockr.ly/2wqN5Nn>

// Keep containers alive during daemon downtime: <http://dockr.ly/2emLwb5>

// Isolate containers with a user namespace: <http://dockr.ly/2gmyKdf>

// Intro to Windows Containers: <https://dockr.ly/2CTYhYb>

CREATING IMAGES

# DISCUSSION: PROVISIONING FILESYSTEMS

What are some potential difficulties with provisioning entire filesystems for containers? How can we avoid these problems?

# LEARNING OBJECTIVES

By the end of this module, learners will be able to

// Create images via several methods

// Describe the filesystem structure underlying an image

// Understand the performance implications of different image design decisions

// Correctly tag and namespace images for distribution on a registry

# WHAT ARE IMAGES?

// A **filesystem** for container process  
// Made of a stack of **immutable** layers  
// Start with a **base image**  
// New layer for each change

Image =  
layered FS

3. Add static files



2. Add EAP



1. RHEL

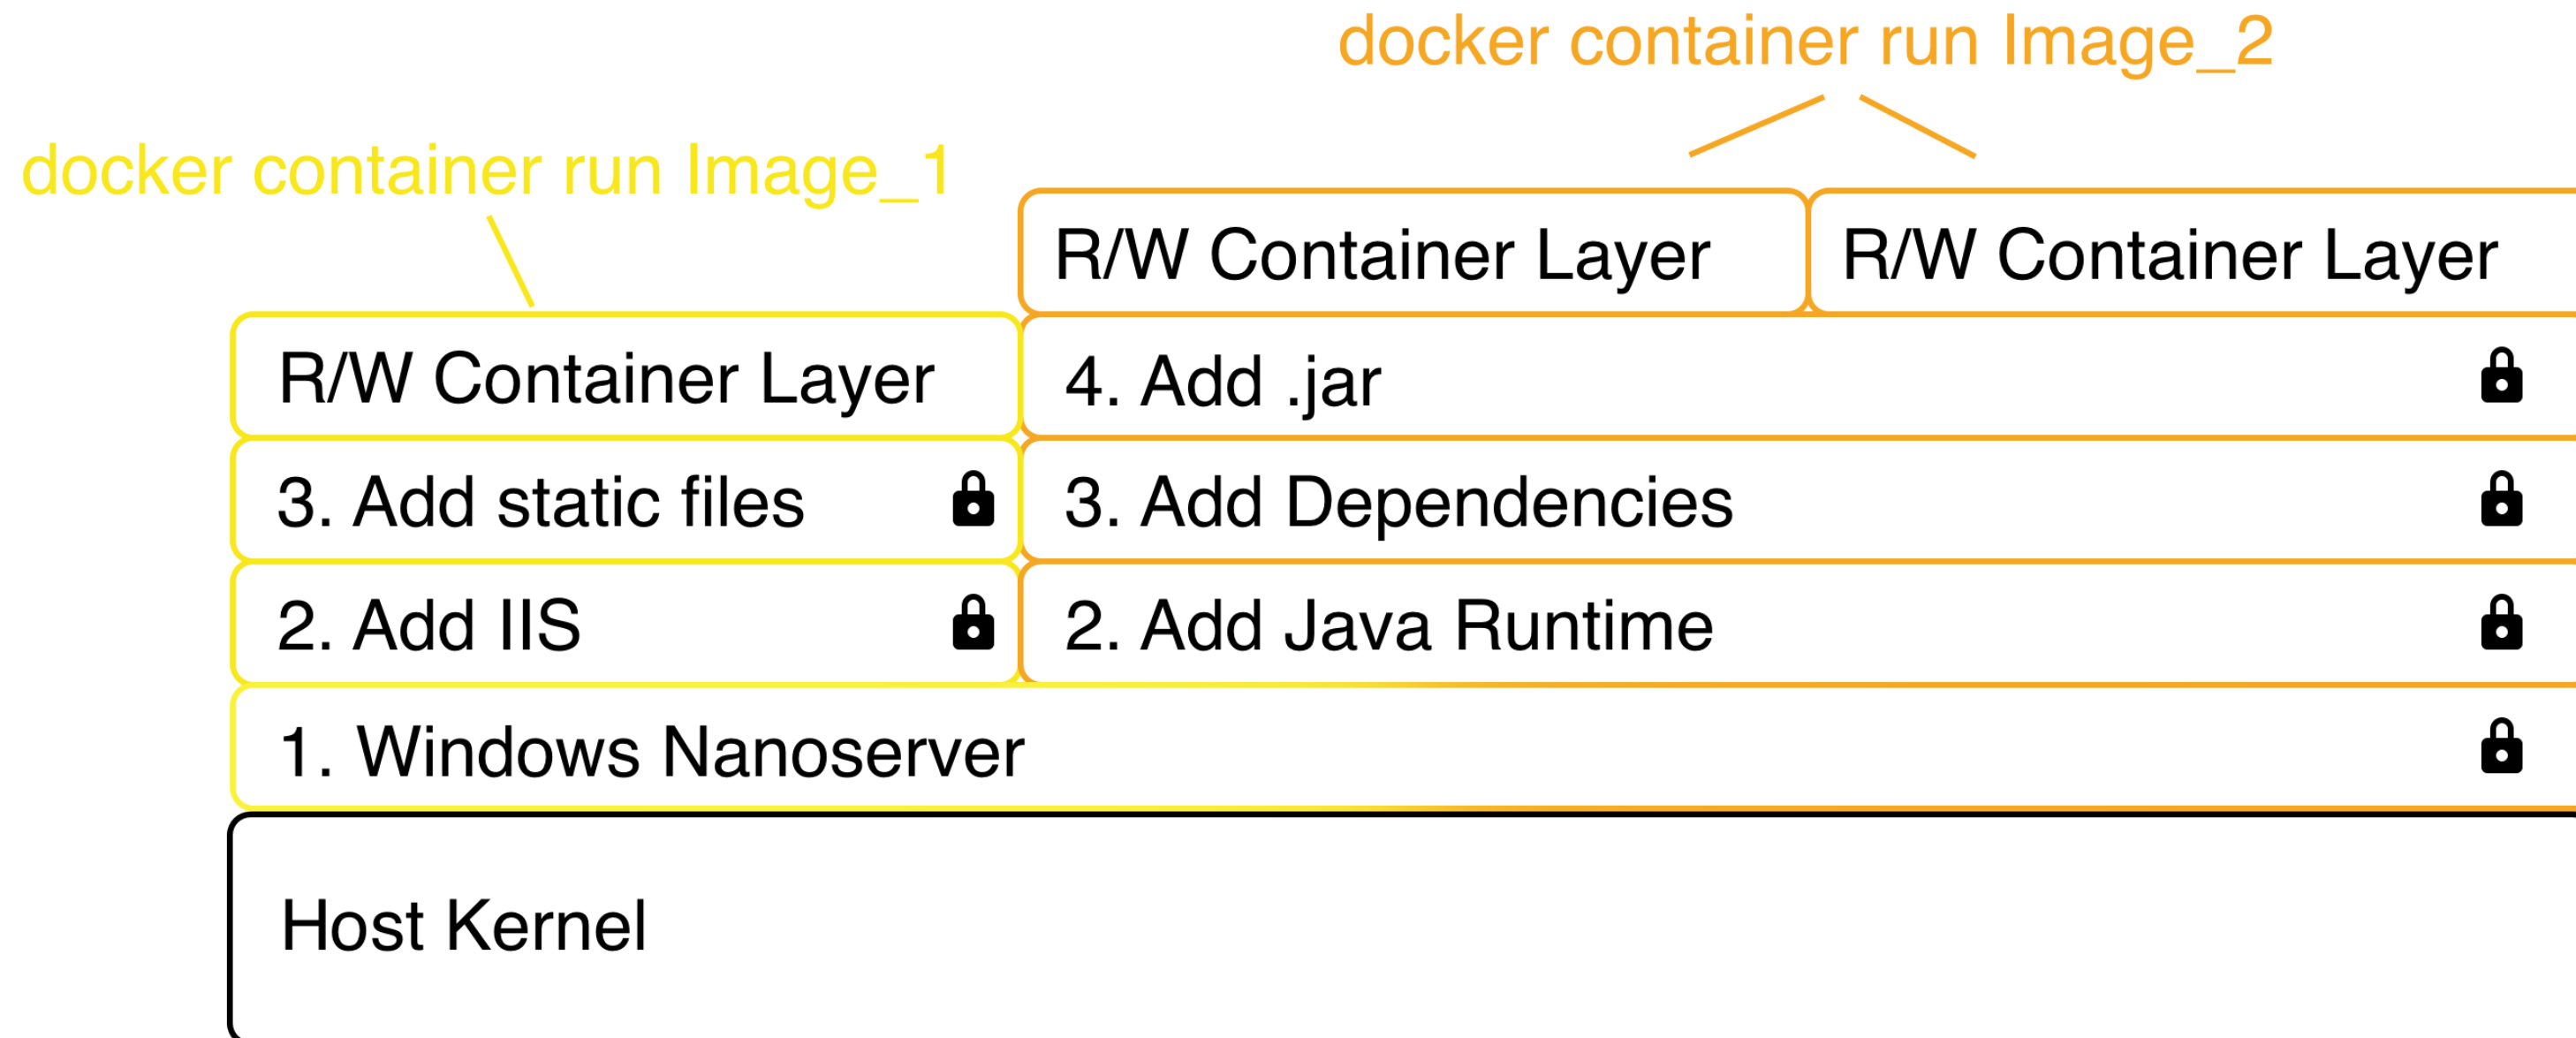




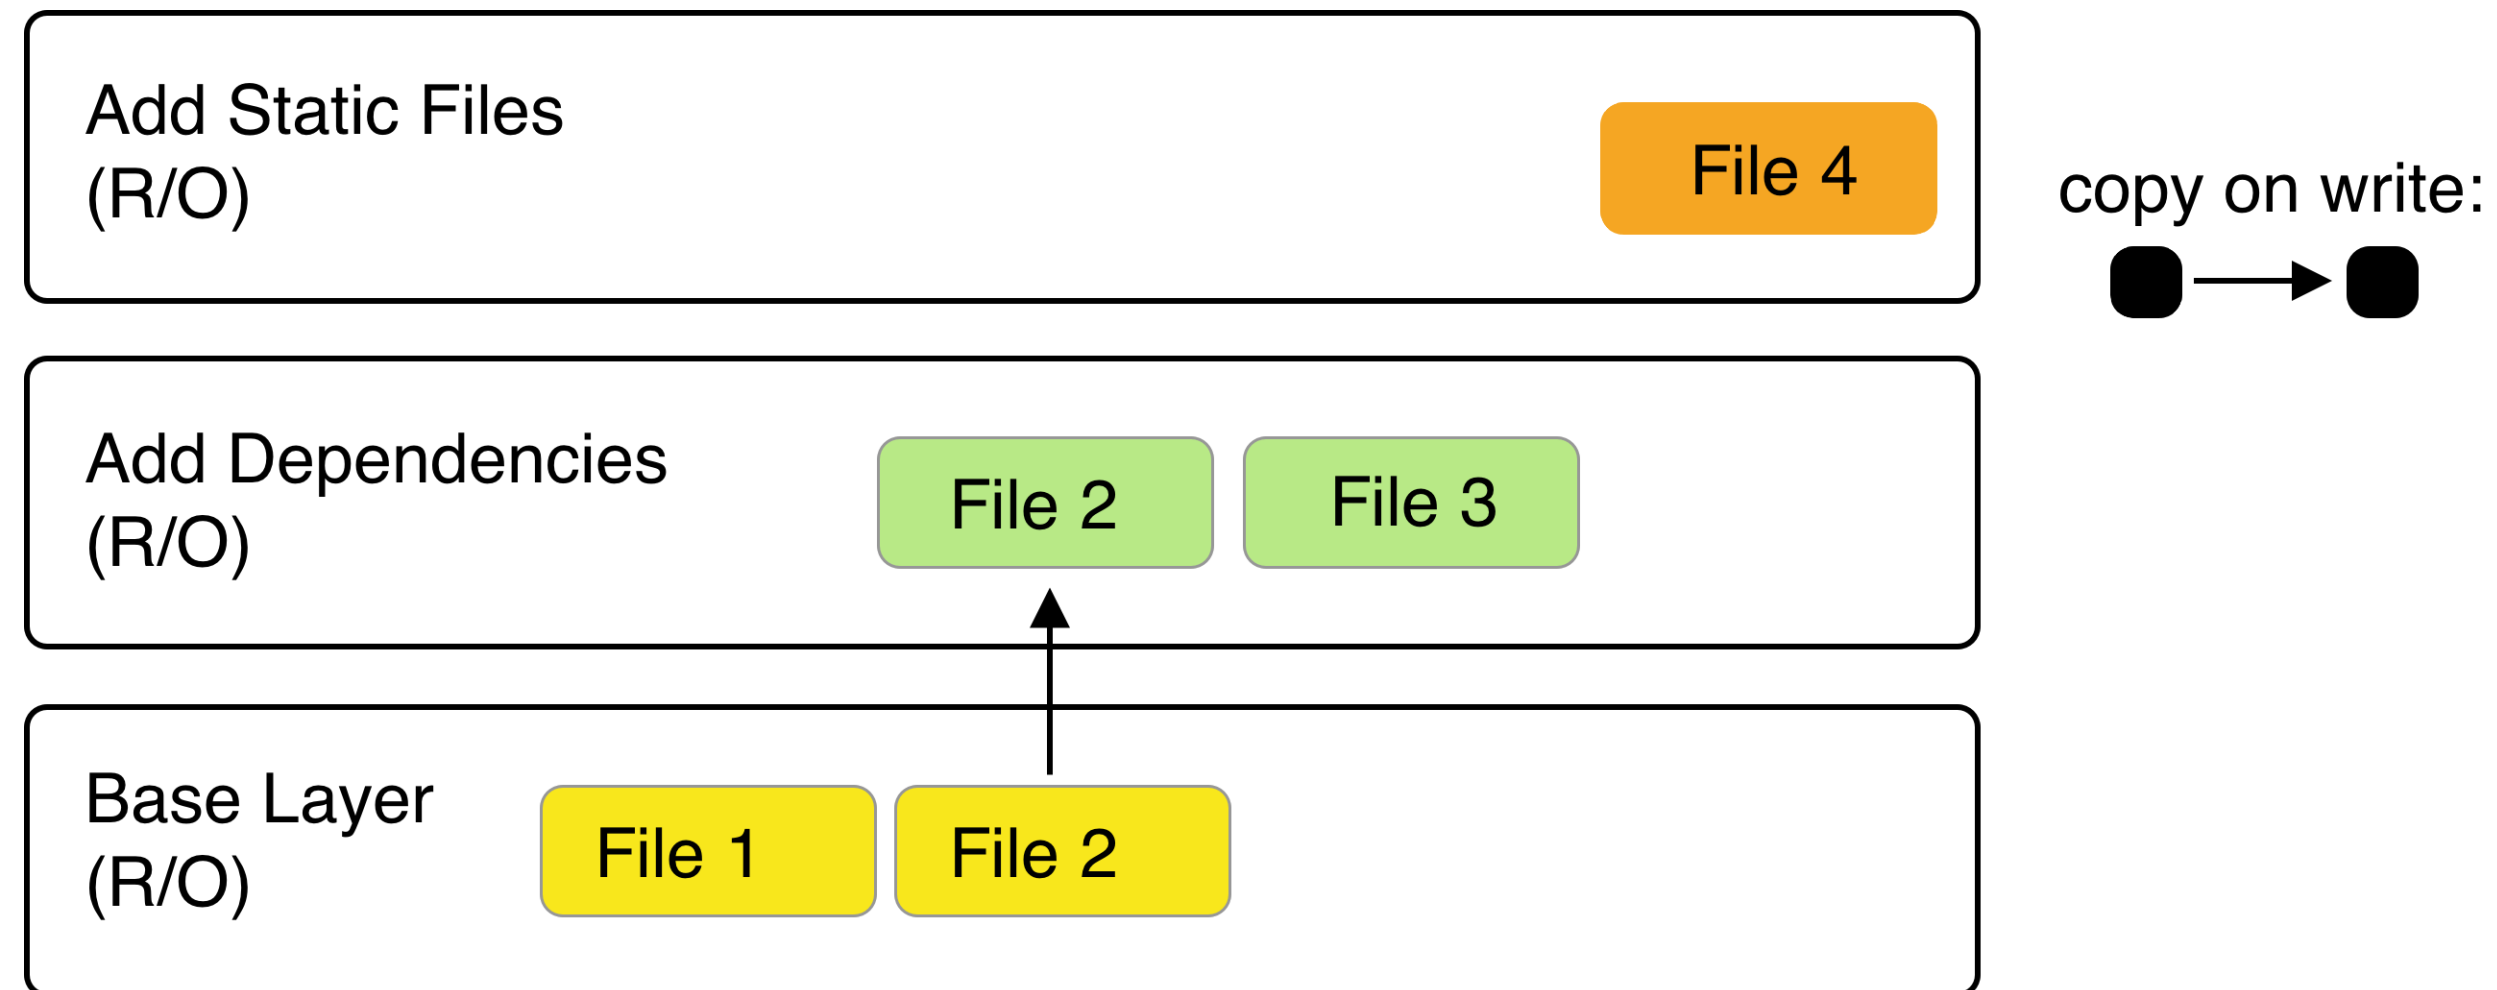
# SHARING LAYERS



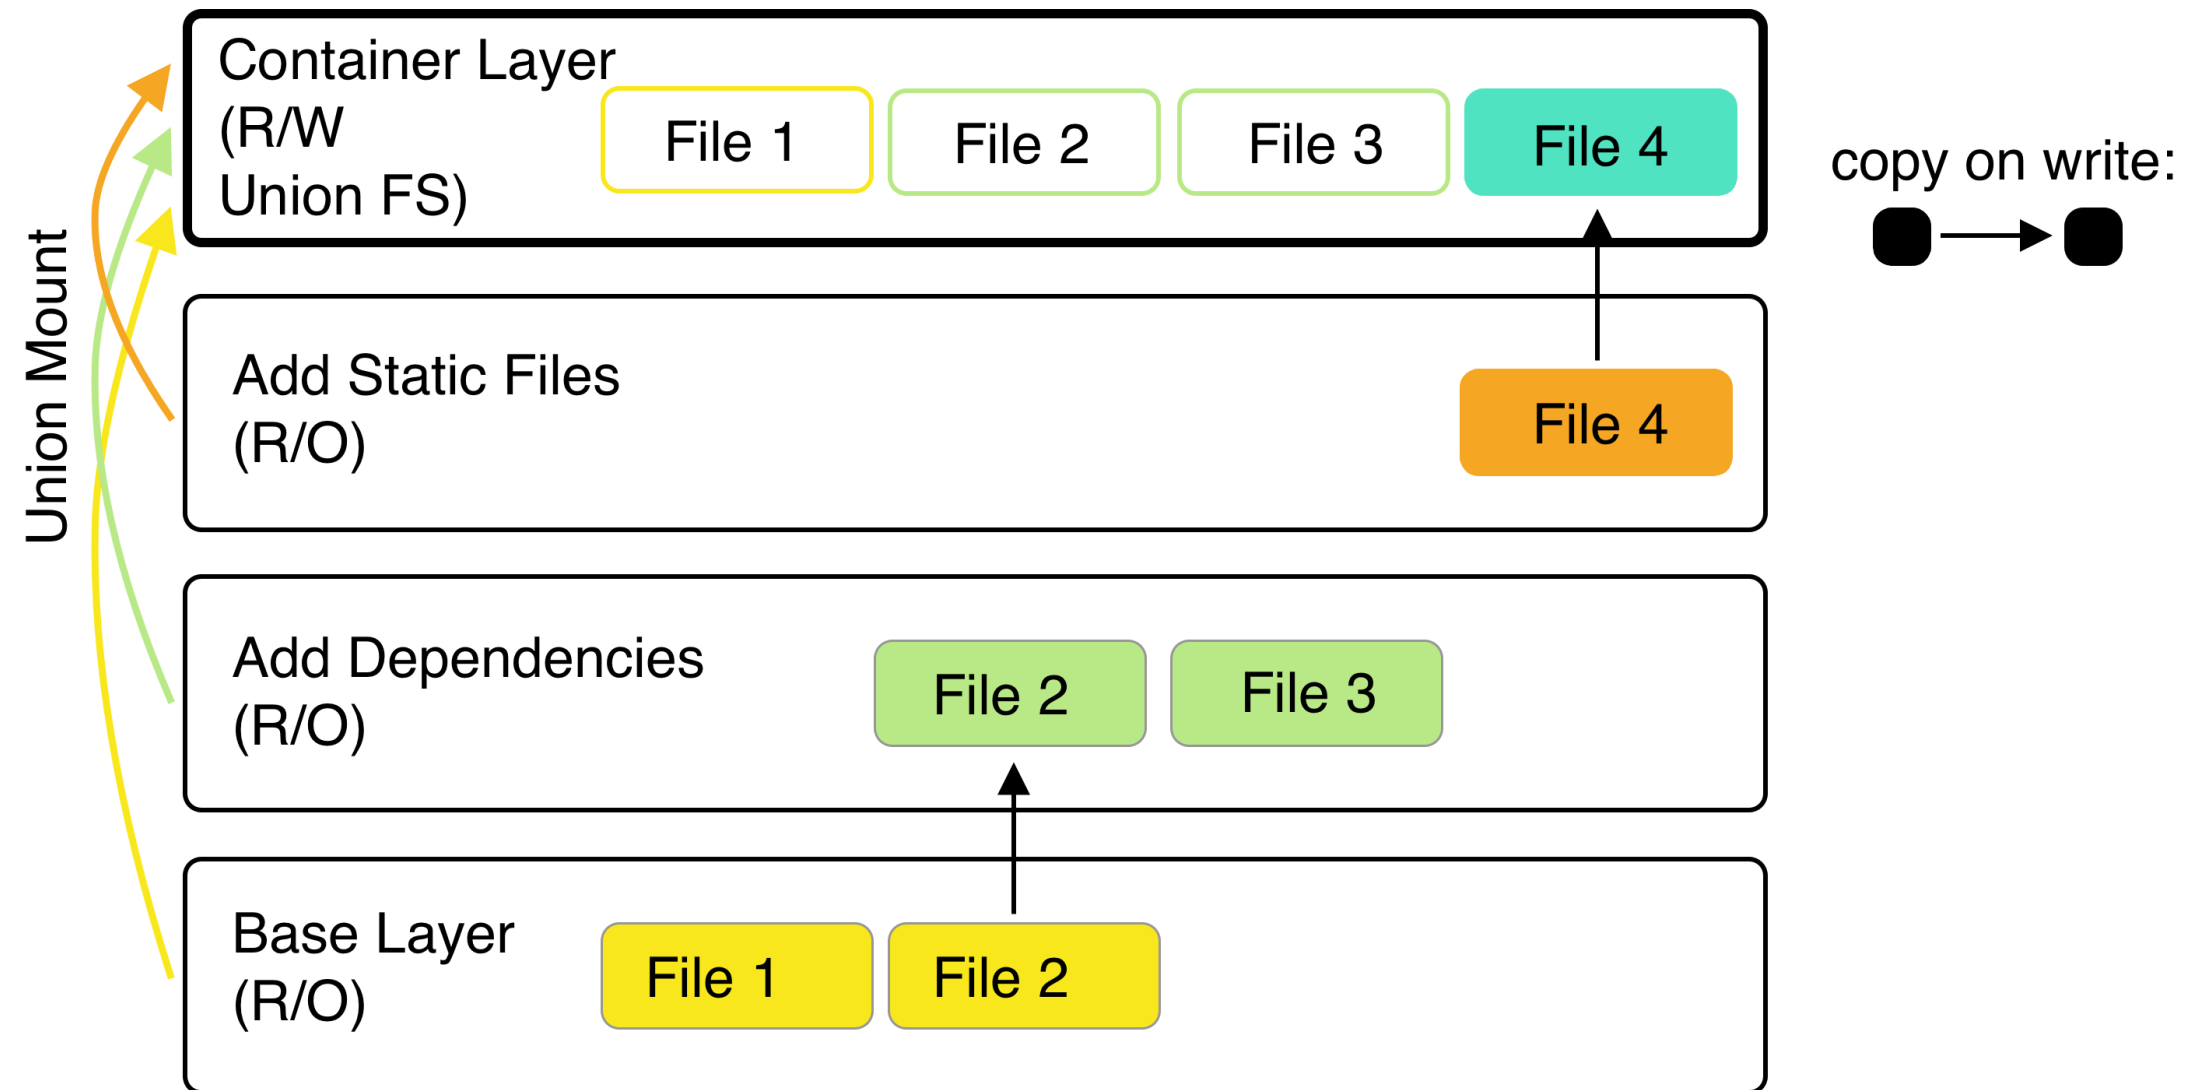
# THE WRITABLE CONTAINER LAYER



# IMAGES: COPY ON WRITE



# LINUX CONTAINERS: UNION FS



# CREATING IMAGES

Three methods:

- // **Commit** the R/W container layer as a new R/O image layer.
- // Define new layers to add to a starting image in a **Dockerfile**.
- // **Import** a tarball into Docker as a standalone base layer.

# COMMITTING CONTAINER CHANGES

```
// docker container commit
```

saves container layer as new R/O image layer

```
// Pro: build images interactively
```

```
// Con: hard to reproduce or audit; avoid this in practice.
```

# DOCKERFILES

// Content manifest

// Provides image layer documentation

// Enables automation (CI/CD)

# DOCKERFILES

// FROM command defines base image.

// Each subsequent command adds a layer or metadata

// docker image build ... builds image from Dockerfile

```
# Comments begin with the pound sign
FROM ubuntu:16.04
RUN apt-get update && apt-get install -y wget
ADD /data /myapp/data
...
```





# EXERCISES: CREATING IMAGES

Work through

// Interactive Image Creation

// Creating Images with Dockerfiles (1/2)

in the Exercises book.

20M 28S




# INSTRUCTOR DEMO: CREATING IMAGES

See the demo

// Creating Images

In the Exercises book.

# BUILD CACHE



// After completion, the resulting image layer is labeled with a hash of the content of all current image layers in the stack.

# CMD AND ENTRYPOINT

// Recall all containers run a process as their PID 1

// `CMD` and `ENTRYPOINT` allow us to specify default processes.

# CMD AND ENTRYPOINT

// `CMD` alone: default command and list of parameters.

// `CMD` + `ENTRYPOINT`: `ENTRYPOINT` provides command, `CMD` provides default parameters.

// `CMD` overridden by command arguments to `docker container run`

// `ENTRYPOINT` overridden via

`--entrypoint` flag to `docker container run`.

# SHELL VS. EXEC FORMAT

```
# Shell form
CMD sudo -u ${USER} java ...

# Exec form
CMD ["sudo", "-u", "jdoe", "java", ...]
```



## EXERCISE: DOCKERFILES (2/2)

Work through

// Creating Images with Dockerfiles (2/2)

In the Exercises book.

20M 14S

# COPY AND ADD COMMANDS

**COPY** copies files from build context to image

```
COPY <src> <dest>
```

**ADD** can also **untar**\* or **fetch URLs**.

*\*Linux containers only!*

In both cases

// create checksum for files added

// log checksum in build cache

// cache invalidated if checksum changed



# DOCKERFILE COMMAND ROUNDUP

// **FROM**: base image to start from (usually OS)

// **RUN**: run a command in the environment defined so far

// **CMD** and **ENTRYPOINT**: define default behavior

// **COPY** and **ADD**: copy files into container

Many more Dockerfile commands are available; see the docs at

<https://docs.docker.com/engine/reference/builder/>

# ADVANCED DOCKERFILE CONSTRUCTION

How can we build images that are

// Lightweight

// Secure

// Minimal build times

# THE SCRATCH IMAGE

// An "empty" image

// Can't be pulled

// Doesn't create a layer

// Used for building images not based on any pre-existing image

// Linux only

```
FROM scratch

ADD centos-7-docker.tar.xz /

LABEL org.label-schema.schema-version="1.0" \
org.label-schema.name="CentOS Base Image" \
org.label-schema.vendor="CentOS" \
org.label-schema.license="GPLv2" \
org.label-schema.build-date="20181205"

CMD ["/bin/bash"]
```

# MULTI-STAGE BUILDS

Hello World, in C:

```
FROM alpine:3.5
RUN apk update && \
    apk add --update alpine-sdk
RUN mkdir /app
WORKDIR /app
ADD hello.c /app
RUN mkdir bin
RUN gcc -Wall hello.c -o bin/hello
CMD /app/bin/hello
```

Builds to:

\$ docker image ls hwc				
REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
hwc	latest	142c29686b6a	15 hours ago	184 MB

# MULTI-STAGE BUILDS

Hello World, lightweight:

```
# Full SDK version (built and discarded)
FROM alpine:3.5 AS build
RUN apk update && \
    apk add --update alpine-sdk
RUN mkdir /app
WORKDIR /app
ADD hello.c /app
RUN mkdir bin
RUN gcc -Wall hello.c -o bin/hello

# Lightweight image returned as final product
FROM alpine:3.5
COPY --from=build /app/bin/hello /app/hello
CMD /app/hello
```

Builds to:

```
$ docker image ls hwc
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
hwc	latest	5d925cfc9c96	39 seconds ago	4MB

# BUILD TARGETS

## Dockerfile

```
FROM <base image> as base
...

FROM <foo image> as foo
...

FROM <bar image> as bar
...

FROM alpine:3.4
...
COPY --from foo ...
COPY --from bar ...
...
```

## Building the image

```
docker image build --target <name> ...
```



# EXERCISE: MULTI-STAGE BUILDS

Work through

// Multi-Stage Builds

In the Exercises book.

19M 57S

# IMAGE CONSTRUCTION BEST PRACTICES

// Start with official images

// Use multi-stage builds to drop compilers, SDKs...

// More layers leverage the cache...

// ...but fewer layers perform better.



# DEVELOPMENT: MORE LAYERS

Bad caching:

```
FROM python:3.5-alpine
RUN mkdir /app
COPY /mypy /app/
RUN pip install -r app/reqs.txt
...
```

Good caching:

```
FROM python:3.5-alpine
RUN mkdir /app
COPY /mypy/reqs.txt /app/
RUN pip install -r app/reqs.txt
COPY /mypy /app/
...
```

# PRODUCTION: LESS LAYERS

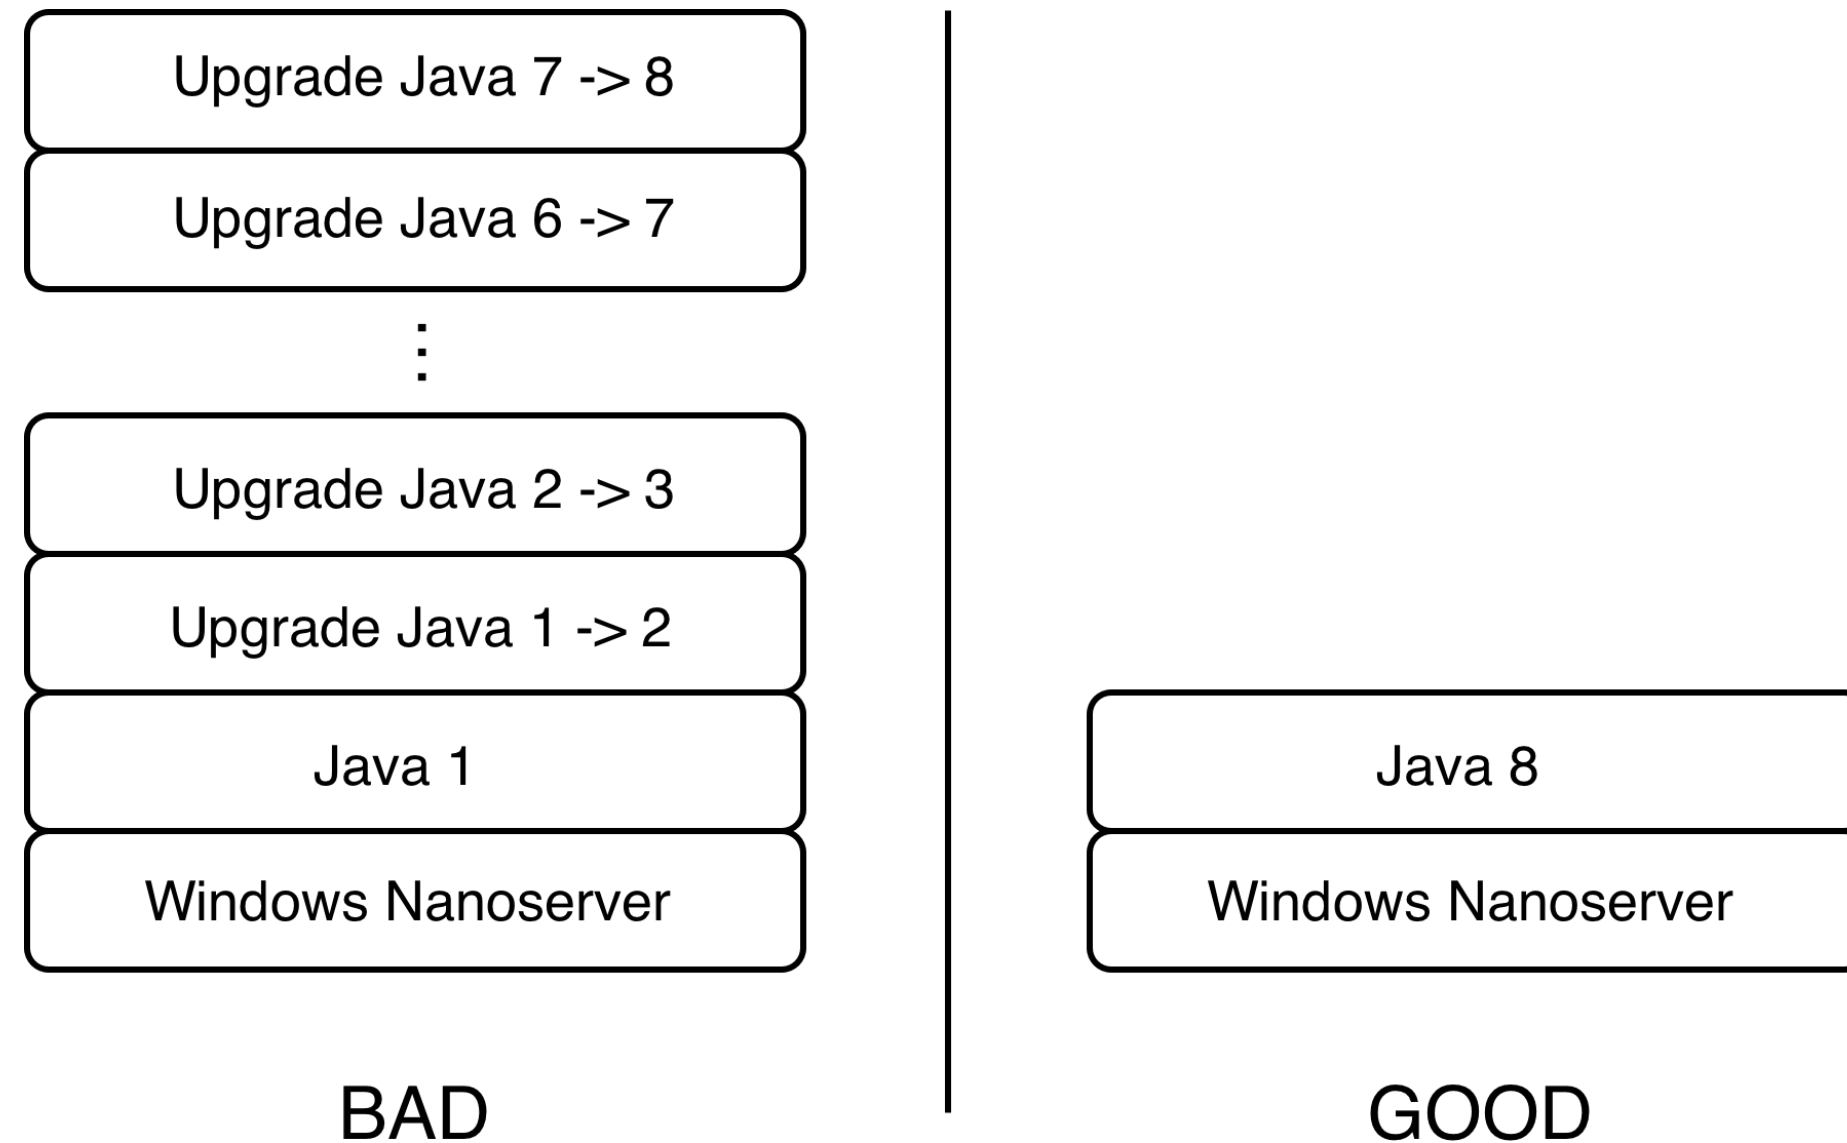
// To collapse ALL image layers:

```
docker container run -d --name demo mytallimage:1.0
docker container export demo > image.tar
cat image.tar | docker image import - myflatimage:1.0
```

// Or build with `--squash` flag (experimental): compress all non-base layers

// Combine `container export` with `--squash` for one shareable base layer + one application-specific upper layer

# BEST PRACTICE: PATCHING & UPDATES



# IMAGE TAGS

// Optional string after image name, separated by :

// :latest by default

// Same image with two tags shares same ID, image layers:

```
$ docker image ls centos*
REPOSITORY      TAG                IMAGE ID           CREATED            SIZE
centos           7                 8140d0c64310      7 days ago        193 MB
$ docker image tag centos:7 centos:mytag
$ docker image ls centos*
REPOSITORY      TAG                IMAGE ID           CREATED            SIZE
centos           7                 8140d0c64310      7 days ago        193 MB
centos           mytag             8140d0c64310      7 days ago        193 MB
```

# IMAGE NAMESPACES

Images exist in one of three namespaces:

```
// Root(ubuntu, nginx, mongo, mysql, ...)
```

```
// User / Org(jdoe/myapp:1.1, microsoft/nanoserver:latest, ...)
```

```
// Registry(FQDN/jdoe/myapp:1.1, ...)
```

# IMAGE TAGGING & NAMESPACE

// Tag on build:

```
docker image build -t myapp:1.0 .
```

// Retag an existing image:

```
docker image tag myapp:1.0 me/myapp:2.0
```

// Note `docker image tag` can set both tag and namespace.

// Names and tags are just pointers to image ID

// Image ID corresponds to immutable content addressable storage

# SHARING IMAGES

// Docker Hub

- Provides certified commercial and free software distributed as Docker Images
- Shares community-generated images and content



# EXERCISE: MANAGING IMAGES

Work through

// Managing Images

In the Exercises book.

19M 37S



# IMAGE CREATION TAKEAWAYS

// Images are built out of read-only layers.

// Dockerfiles specify image layer contents.

// Key Dockerfile commands: `FROM`, `RUN`, `COPY` and `ENTRYPOINT`

// Images must be namespaced according to where you intend on sharing them.

# FURTHER READING

// Best practices for writing Dockerfiles: <http://dockr.ly/22WiJi0>

// Use multi-stage builds: <http://dockr.ly/2ewcUY3>

// More about images, containers, and storage drivers: <http://dockr.ly/1TuWndC>

// Details on image layering: <https://bit.ly/2AHX7iW>

// Graphdriver plugins: <http://dockr.ly/2eIVCab>

// Docker Reference Architecture: An Intro to Storage Solutions for Docker CaaS: <http://dockr.ly/2x8sBw2>

// How to select a storage driver: <http://dockr.ly/2eDu8y0>

// Use the AUFS storage driver: <http://dockr.ly/2jVc1Zz>

// User guided caching in Docker: <http://dockr.ly/2xKafPf>

# DOCKER VOLUMES

# DISCUSSION: MANAGING DATA

If a container generates a lot of data, where should it be stored? What if you need to provision data to a container?

# LEARNING OBJECTIVES

By the end of this module, learners will be able to

// Define a volume and identify its primary use cases

// Describe the advantages and potential security risks of mounting volumes and host directories into containers

# VOLUME USECASES

Volumes provide a R/W path **separate from the layered filesystem.**

// **Mount** data at container startup

// **Persist** data when a container is deleted

// **Share** data between containers

// **Speed up** I/O by circumventing the union filesystem

# BASIC VOLUMES

// **Named**: managed by Docker; filesystem independent; user-specified identifier

// **Anonymous**: managed by Docker; filesystem independent; randomly-generated identifier

// **Host mounted**: mount a specific path on the host; DIY management



# INSTRUCTOR DEMO: VOLUMES

See the demo

// Basic Volume Usage

In the Exercises book.



# VOLUMES IN DOCKERFILES

// VOLUME instruction creates a mount point

// Can specify arguments in a JSON array or string

// Cannot map volumes to host directories

// Volumes are initialized when the container is executed

```
FROM nginx:latest
...
# string example
VOLUME /myvolume

# string example with multiple volumes
VOLUME /www/website1 /www/website2

# JSON example
VOLUME ["myvol1", "myvol2"]
...
```

# VOLUMES AND SECURITY

// Point of ingress to the host and other containers

// Don't mount things unnecessarily

// Use the `:ro` flag

// Linux: in-memory `tmpfs` mounts available



# EXERCISE: VOLUMES USECASE

Work through

// Database Volumes

In the Exercises book.

19M 15S

# DOCKER VOLUME TAKEAWAYS

// Volumes persist data beyond the container lifecycle

// Volumes bypass the copy on write system (better for write-heavy containers)

# FURTHER READING

// How to use volumes: <http://dockr.ly/2vRZBDG>

// Troubleshoot volume errors: <http://dockr.ly/2vyjvbP>

// Docker volume reference: <http://dockr.ly/2ewrlew>

# DOCKER SYSTEM COMMANDS

# LEARNING OBJECTIVES

By the end of this module, learners will be able to

// Execute clean-up commands

// Locate Docker system information

# CLEAN-UP COMMANDS

```
// docker system df
```

TYPE	TOTAL	ACTIVE	SIZE	RECLAIMABLE
Images	39	2	9.01 GB	7.269 GB (80%)
Containers	2	2	69.36 MB	0 B (0%)
Local Volumes	0	0	0 B	0 B

```
// docker system prune
```

more limited...

```
// docker image prune [--filter "foo=bar"]
```

```
// docker container prune [--filter "foo=bar"]
```

```
// docker volume prune [--filter "foo=bar"]
```

```
// docker network prune [--filter "foo=bar"]
```





# INSPECT THE SYSTEM

## docker system info

```
Containers: 2
  Running: 2
  Paused: 0
  Stopped: 0
Images: 105
Server Version: 17.03.0-ee
Storage Driver: overlay2
  Backing Filesystem: extfs
  Supports d_type: true
  Native Overlay Diff: true
Logging Driver: json-file
Cgroup Driver: cgroupfs
Plugins:
  Volume: local
  Network: bridge host ipvlan macvlan null overlay
Swarm: active
  NodeID: ybmqksh6fm627armruq0e8id1
  Is Manager: true
  ClusterID: 2rbf1dv6t5ntro2fxbry6ikr3
  Managers: 1
  Nodes: 1
  Orchestration:
    Task History Retention Limit: 5
  Raft:
    Snapshot Interval: 10000
    Number of Old Snapshots to Retain: 0
    Heartbeat Tick: 1
  ...
```



# SYSTEM EVENTS

Start observing with ...

```
docker system events
```

Generate events with ...

```
docker container run --rm alpine echo 'Hello World!'
```

```
2017-01-25T16:57:48.553596179-06:00 container create 30eb630790d44052f26c1081...
2017-01-25T16:57:48.556718161-06:00 container attach 30eb630790d44052f26c1081...
2017-01-25T16:57:48.698190608-06:00 network connect de1b2b40f522e69318847ada3...
2017-01-25T16:57:49.062631155-06:00 container start 30eb630790d44052f26c1081d...
2017-01-25T16:57:49.164526268-06:00 container die 30eb630790d44052f26c1081dbf...
2017-01-25T16:57:49.613422740-06:00 network disconnect de1b2b40f522e69318847a...
2017-01-25T16:57:49.815845051-06:00 container destroy 30eb630790d44052f26c108...
```





## EXERCISE: SYSTEM COMMANDS

Work through

// Cleaning up Docker Resources

// Inspection Commands

in the Exercises book.

18M 57S

# DISCUSSION

// What is the origin of dangling image layers?

// What are some potential pitfalls to automating system cleanup with prune commands, and how to avoid them?

// Questions?

# FURTHER READING

// System commands reference: <http://dockr.ly/2eMR53i>

# CONTAINERIZATION FUNDAMENTALS

## CONCLUSION: ANY APP, ANYWHERE.

// Containers are isolated processes

// Images provide filesystem for containers

// Volumes persist data

WRAP UP - SPRING BOOT

# EXERCISE INSTRUCTIONS

// Goal: Build a docker image that runs a Java application

// Clone the git repository at <https://github.com/ckaserer/java-helloworld>

// Build the application using `gradle bootJar` and find the resulting fat-jar file at `build/libs` (already done for you).

// You can run the java application locally with `java -jar rest-service-0.0.1-SNAPSHOT.jar`, this starts an application server on port 8080. Check it in your browser via <http://127.0.0.1:8080/greeting> (respectively with your AWS instance public dns).

18M 47S



# SOLUTION HINTS

- // Build the application outside of the container (already done and checked into git).
- // Find a suitable docker base image, that has a JRE installation.
- // Use both the `CMD` and the `ENTRYPOINT` instructions.
- // Run two instances of the container – what do you have to take care of? Make sure both instances are reachable via browser.

18M 44S

# SOLUTION

## Sample Dockerfile

```
FROM anapsix/alpine-java
LABEL MAINTAINER=clemens.kaserer@gepardec.com
WORKDIR /data
EXPOSE 8080
COPY build/libs/rest-service-0.0.1-SNAPSHOT.jar \
    rest-service-0.0.1-SNAPSHOT.jar
CMD ["-jar", "rest-service-0.0.1-SNAPSHOT.jar"]
ENTRYPOINT ["java"]
```

## Solution Commands

```
// docker build -t spring_boot_example .
```

```
// docker run -d -p 80:8080 spring_boot_example
```

# CONTAINERIZATION TRAINING

Please take our feedback survey

Get in touch: [office@gepardec.com](mailto:office@gepardec.com)

<https://www.gepardec.com/trainings>

# DOCKER NETWORKING BASICS

# DISCUSSION: PORTABLE NETWORKS

Network traffic must by definition traverse a network outside its originating container. How can we make inter-container communication as portable and secure as containers themselves?

# LEARNING OBJECTIVES

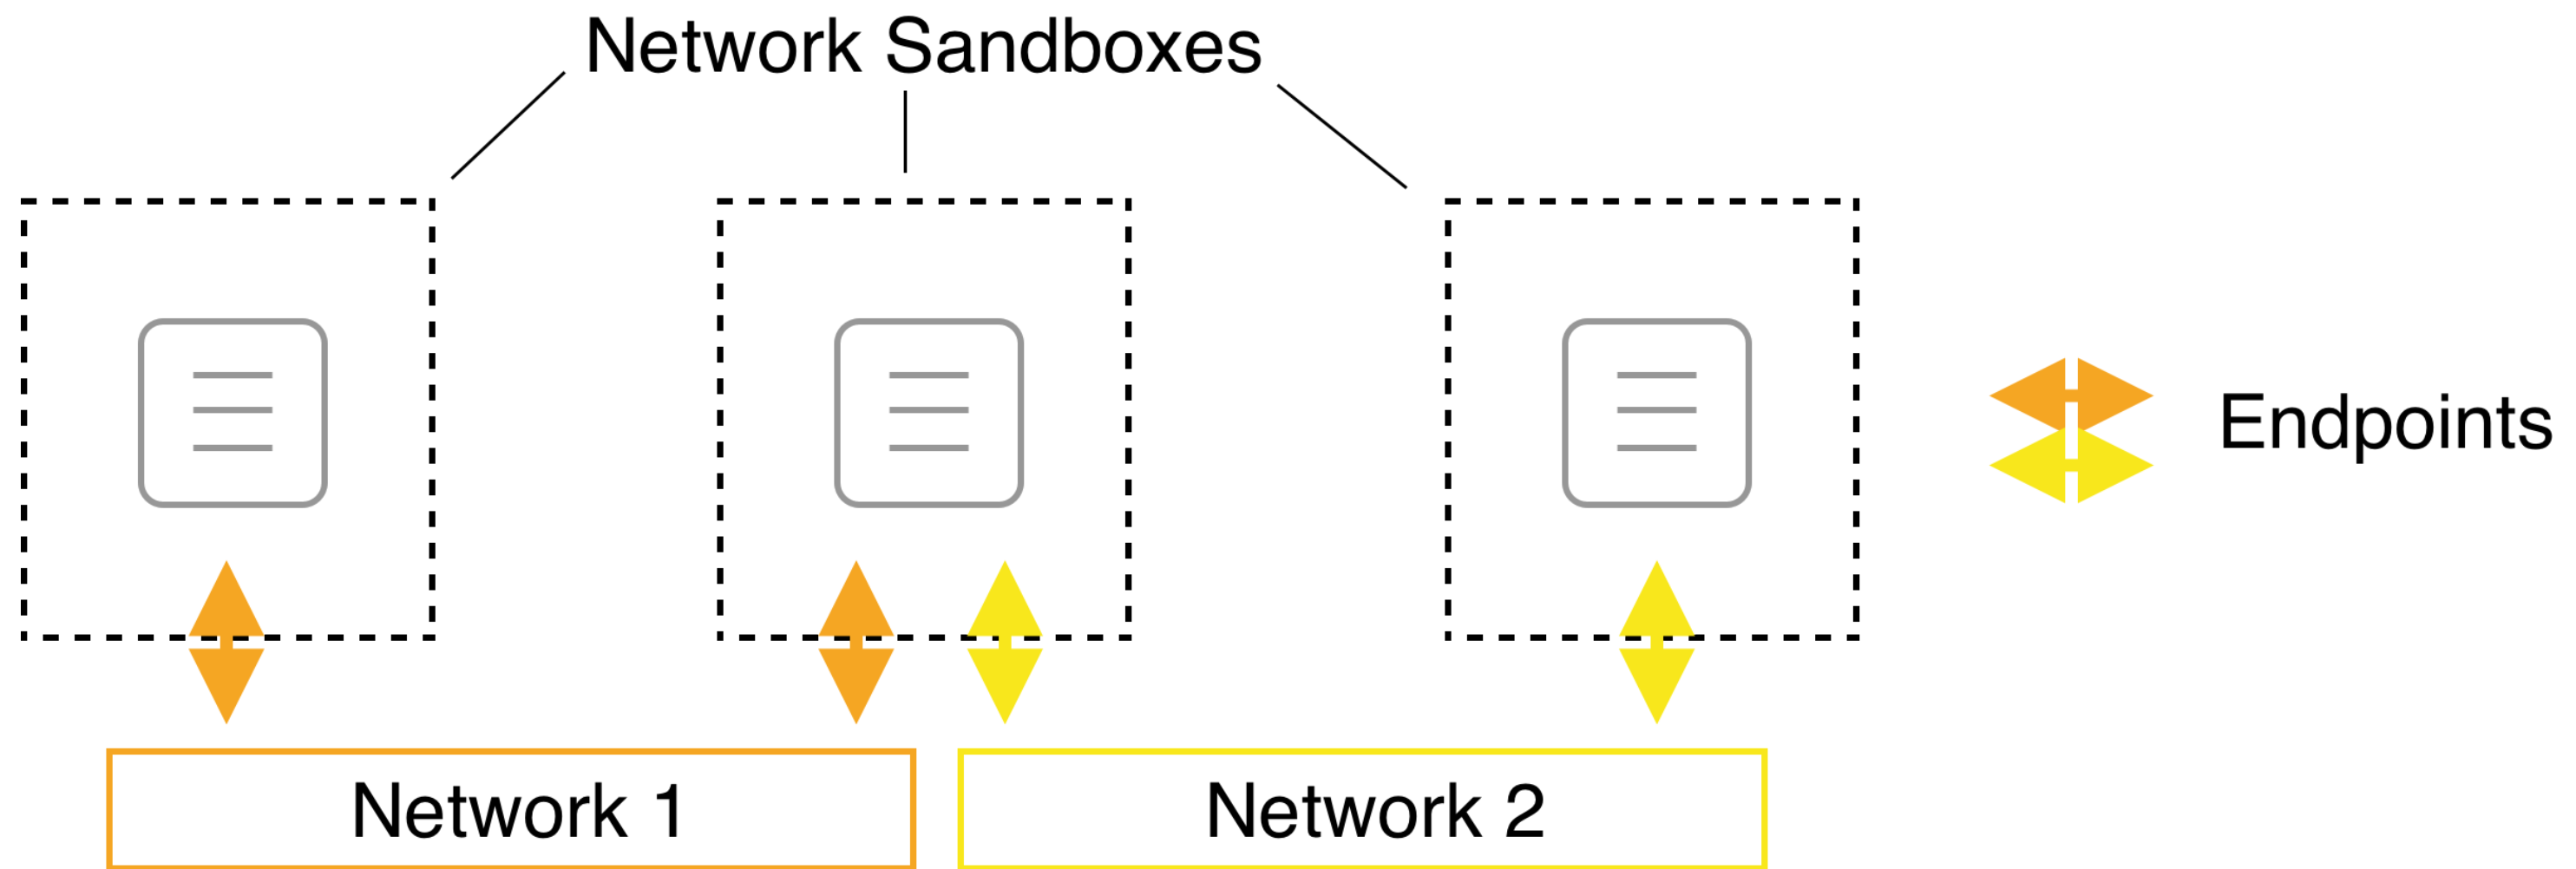
By the end of this module, learners will be able to

// Describe Docker's container network model and its security implications

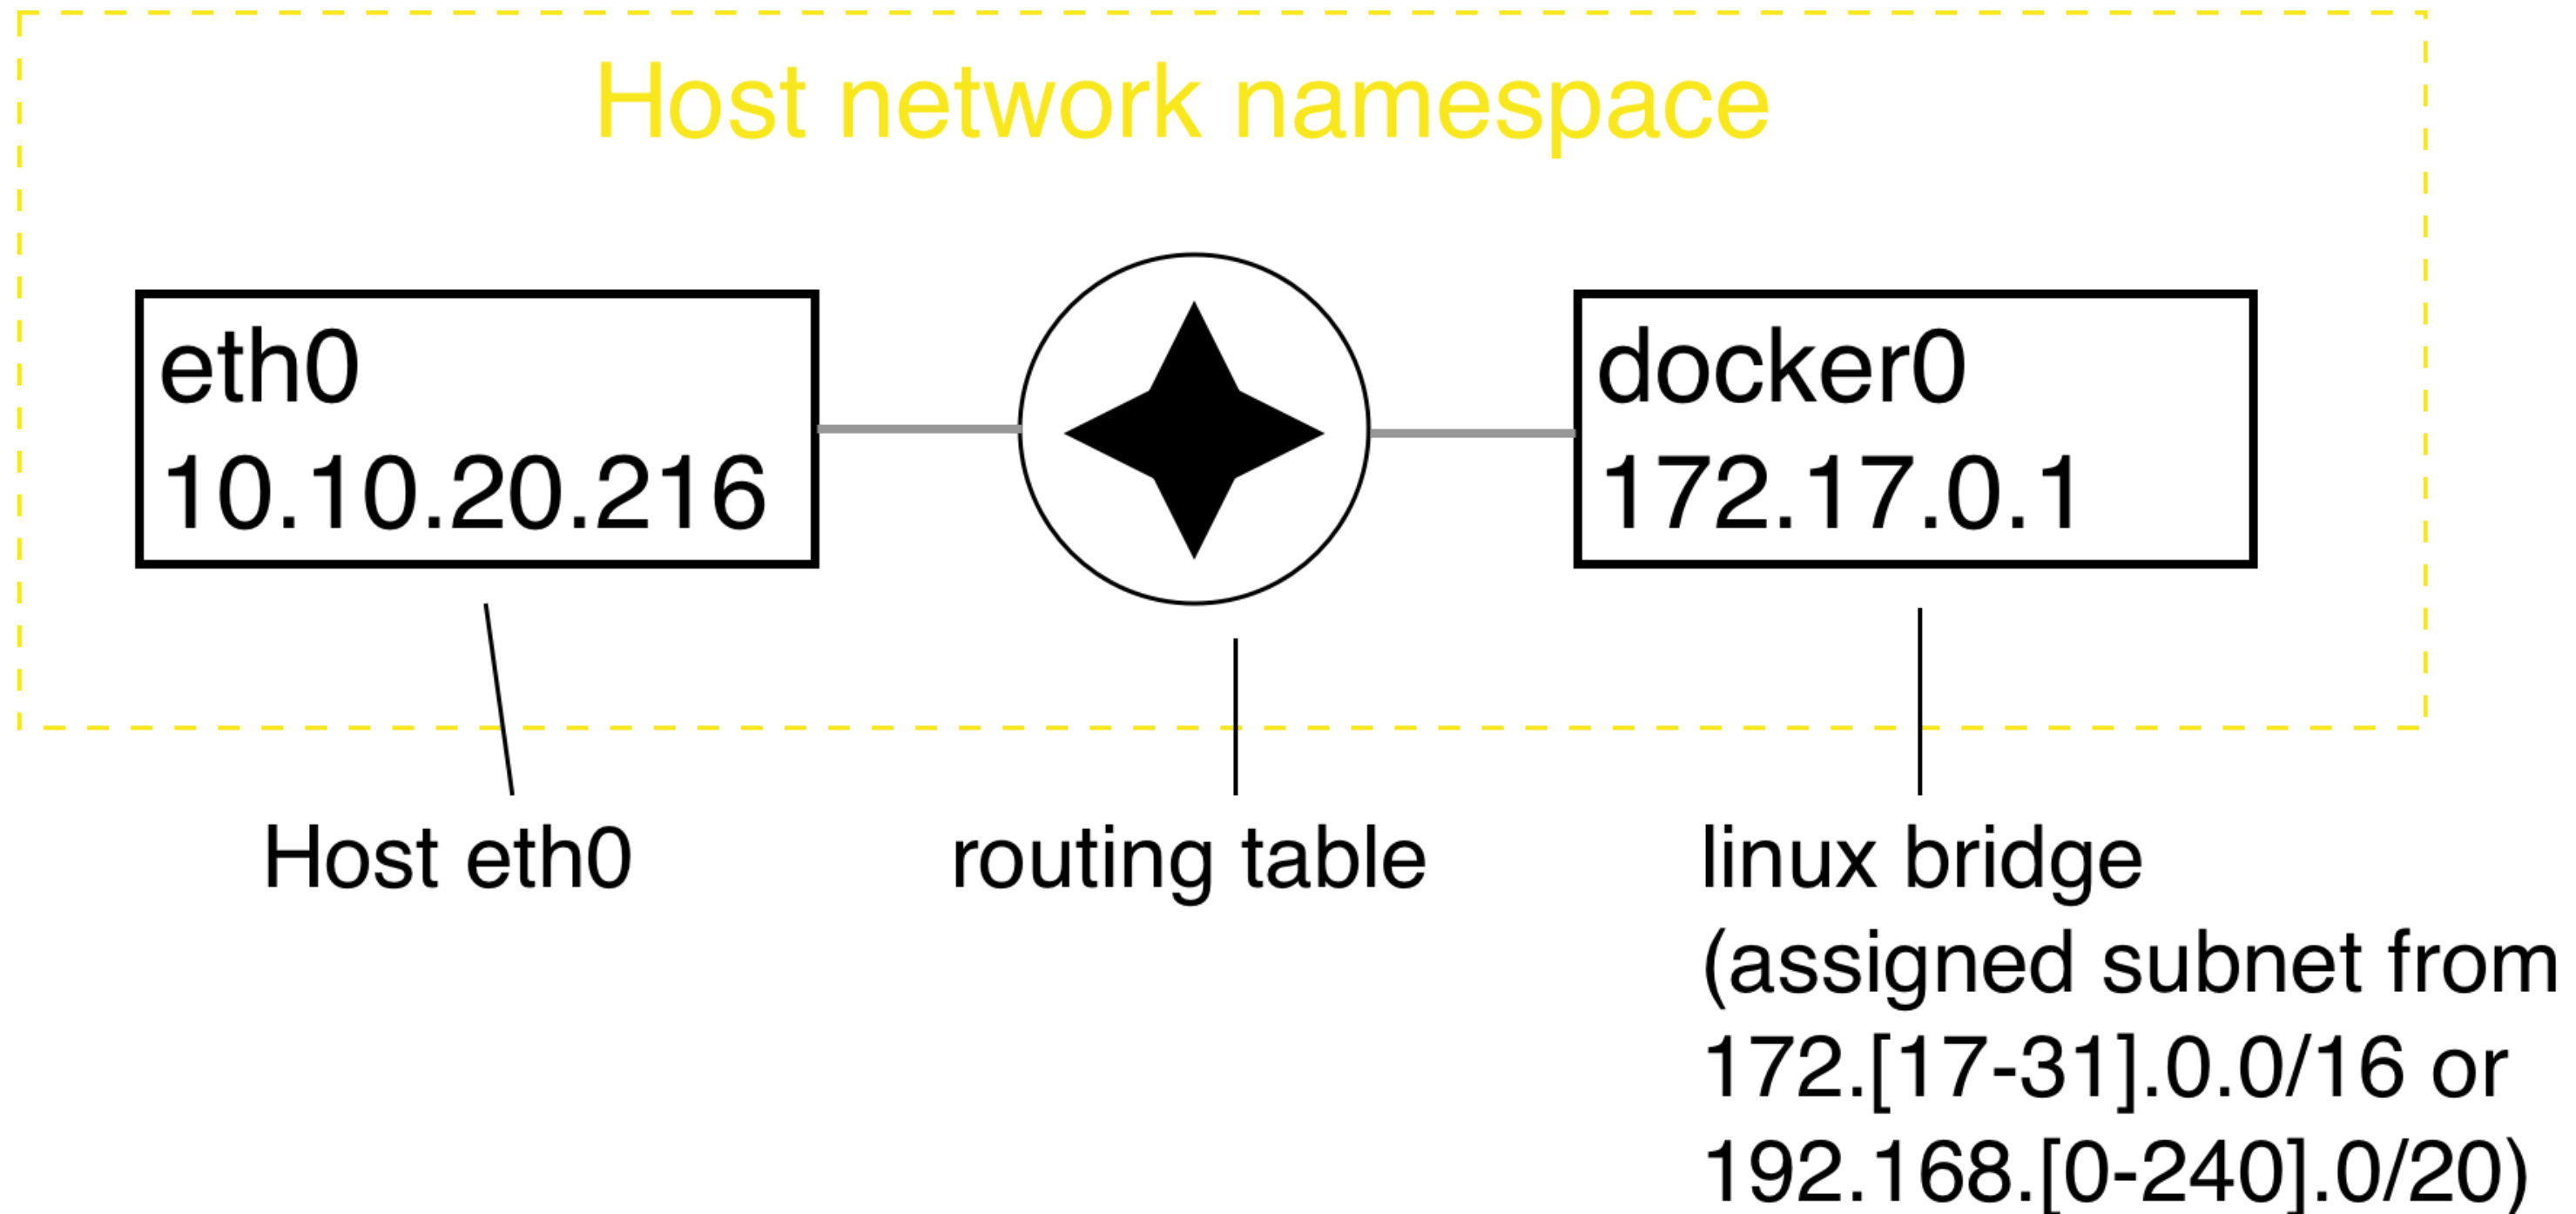
// Describe the basic technologies that underwrite single host networks

// Understand how Docker manipulates a host's firewall rules to control container traffic

# THE CONTAINER NETWORK MODEL

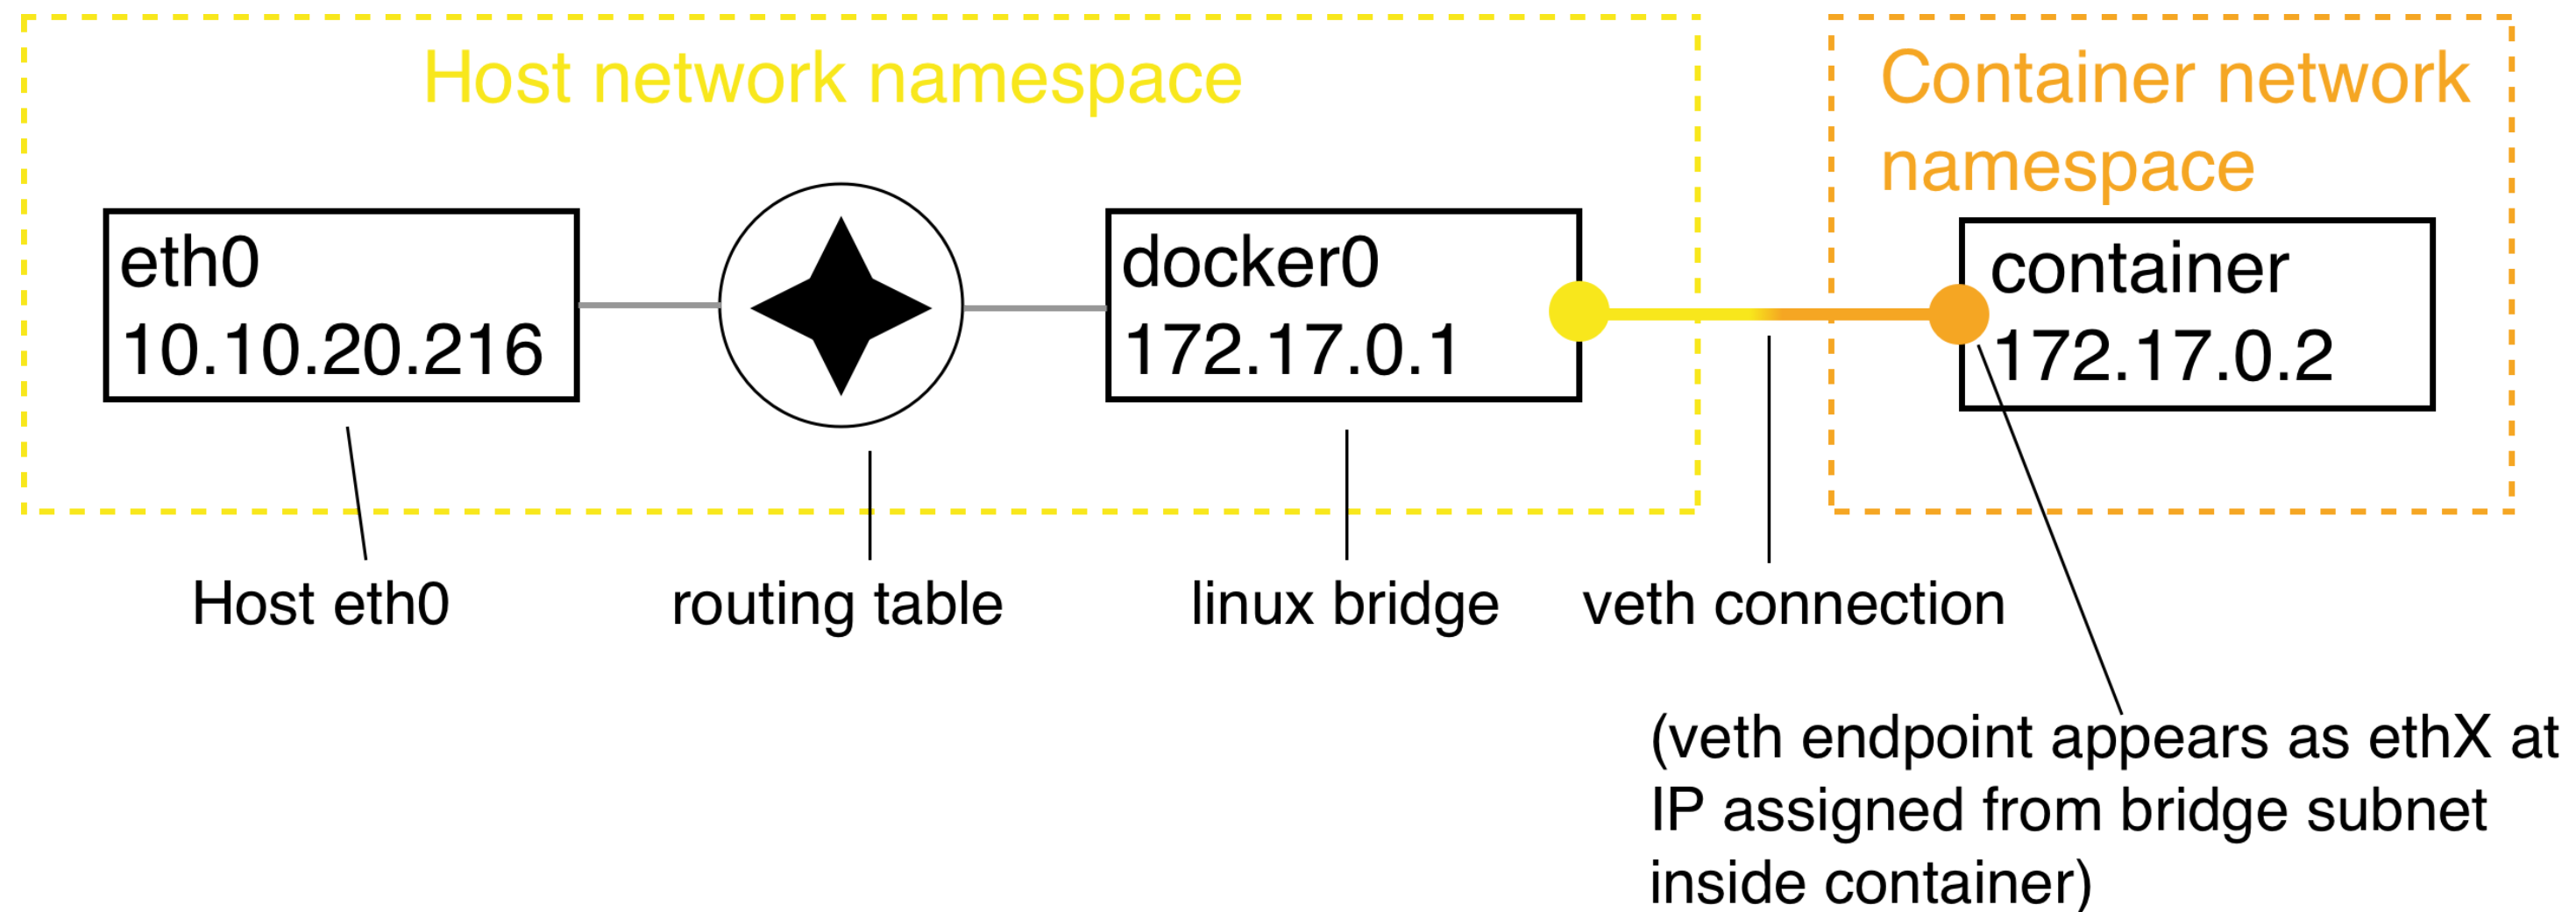


# LINUX: DEFAULT SINGLE-HOST NETWORK



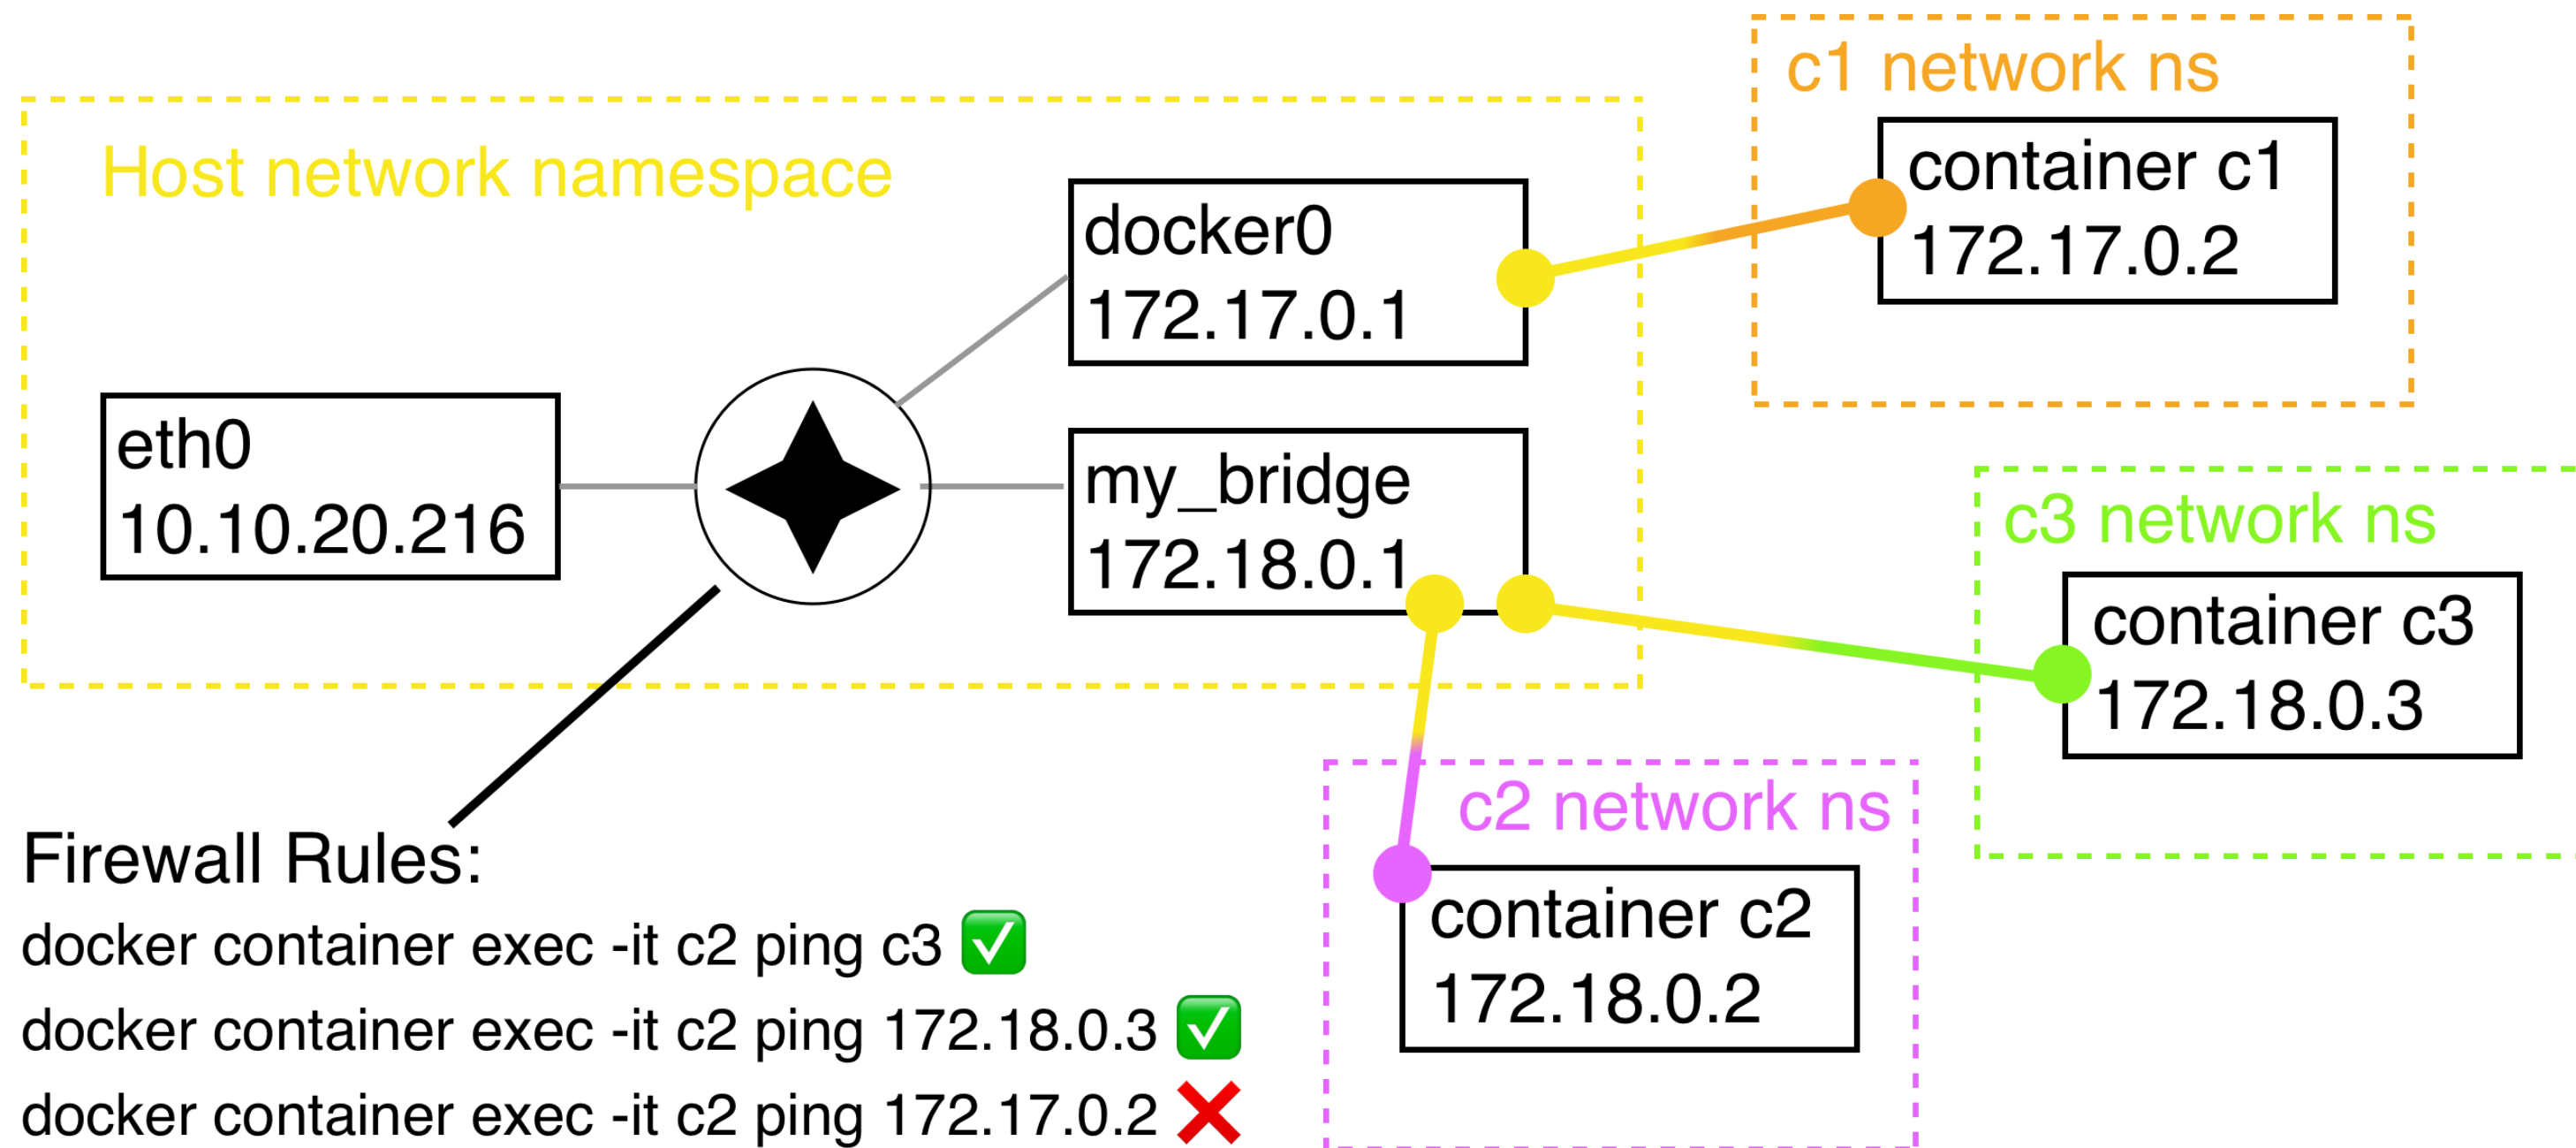


# LINUX: DEFAULT CONTAINER NETWORKING



**Quiz:** identify the sandbox, endpoint and network corresponding to the container networking model objects in this diagram.

# LINUX: USER-DEFINED BRIDGES & FIREWALLS



# EXPOSING CONTAINER PORTS

// Containers have no public IP address by default.

// Can forward host port -> container port

// Mapping created manually or automatically.

// Port mappings visible via

`docker container ls` or

`docker container port`



# INSTRUCTOR DEMO: SINGLE HOST NETWORKS

See the demo

// Single Host Networks

In the Exercises book.



# EXERCISE: SINGLE HOST NETWORKS

Work through

// Introduction to Container Networking

// Container Port Mapping

in the Exercises book.

18M 17S

# DOCKER NETWORKING TAKEAWAYS

// Single host networks follow the container networking model:

- Sandbox: Network namespaces
- Endpoint: veth (linux)
- Network: bridge (linux)

// Containers resolve each other by DNS lookup when explicitly named and attached to custom networks

// Docker software defined networks are firewalled from each other by default

# FURTHER READING

// Docker Reference Architecture: Designing Scalable, Portable Docker Container Networks:

<https://dockr.ly/2q308jq>

// Network containers: <http://dockr.ly/2x1BYgW>

// Docker container networking: <http://dockr.ly/1QnT6y8>

// Understand container communication: <http://dockr.ly/2iSrH00>

# INTRODUCTION TO DOCKER COMPOSE



# DISCUSSION: PROCESSES VS. APPLICATIONS

Containers and images describe individual processes. What will we need to describe entire applications?

# LEARNING OBJECTIVES

By the end of this module, learners will be able to

// Design scalable Docker services

// Leverage Docker's built in service discovery mechanism

// Write a compose file describing an application

# DISTRIBUTED APPLICATION ARCHITECTURE

// Applications consisting of one or more containers across one or more nodes

// Docker Compose facilitates multi-container design **on a single node**

# DOCKER SERVICES

// Goal: declare and (re)configure many similar containers all at once


// Goal: scale apps by adding containers seamlessly

// A **service** defines the **desired state** of a group of identically configured containers


// Docker provides **transparent service discovery** for Services

# SERVICE DISCOVERY

process to process



service to service



Services are assigned a **Virtual IP** which spreads traffic out across the underlying containers automatically.

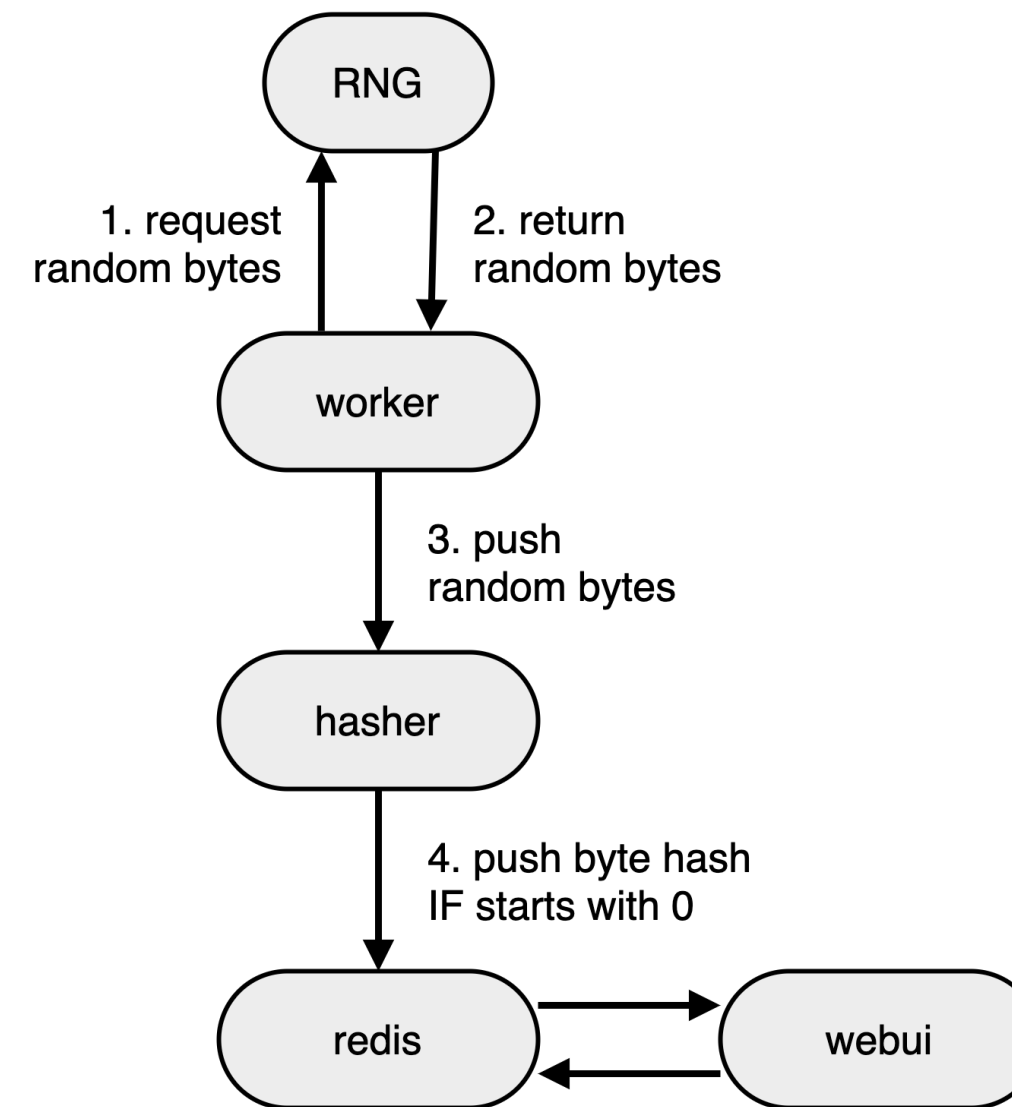
# OUR APPLICATION: DOCKERCOINS



(DockerCoins 2016 logo courtesy of [@XtlCnslt](#) and [@ndeloof](#). Thanks!)

// It is a DockerCoin miner! 💰💣📦🏭

// Dockercoins consists of 5 services working together:





# INSTRUCTOR DEMO: DOCKER COMPOSE

See the demo

// Docker Compose

In the Exercises book.



# EXERCISE: COMPOSE APPS

Work through

// Starting a Compose App

// Scaling a Compose App

in the Exercises book.

17M 54S



# DOCKER COMPOSE TAKEAWAYS

// Docker Compose makes single node orchestration easy

// Compose services makes scaling applications easy

// Bottleneck identification important

// Syntactically: `docker-compose.yml` + API

# FURTHER READING

// Docker compose examples: <http://dockr.ly/1FL2VQ6>

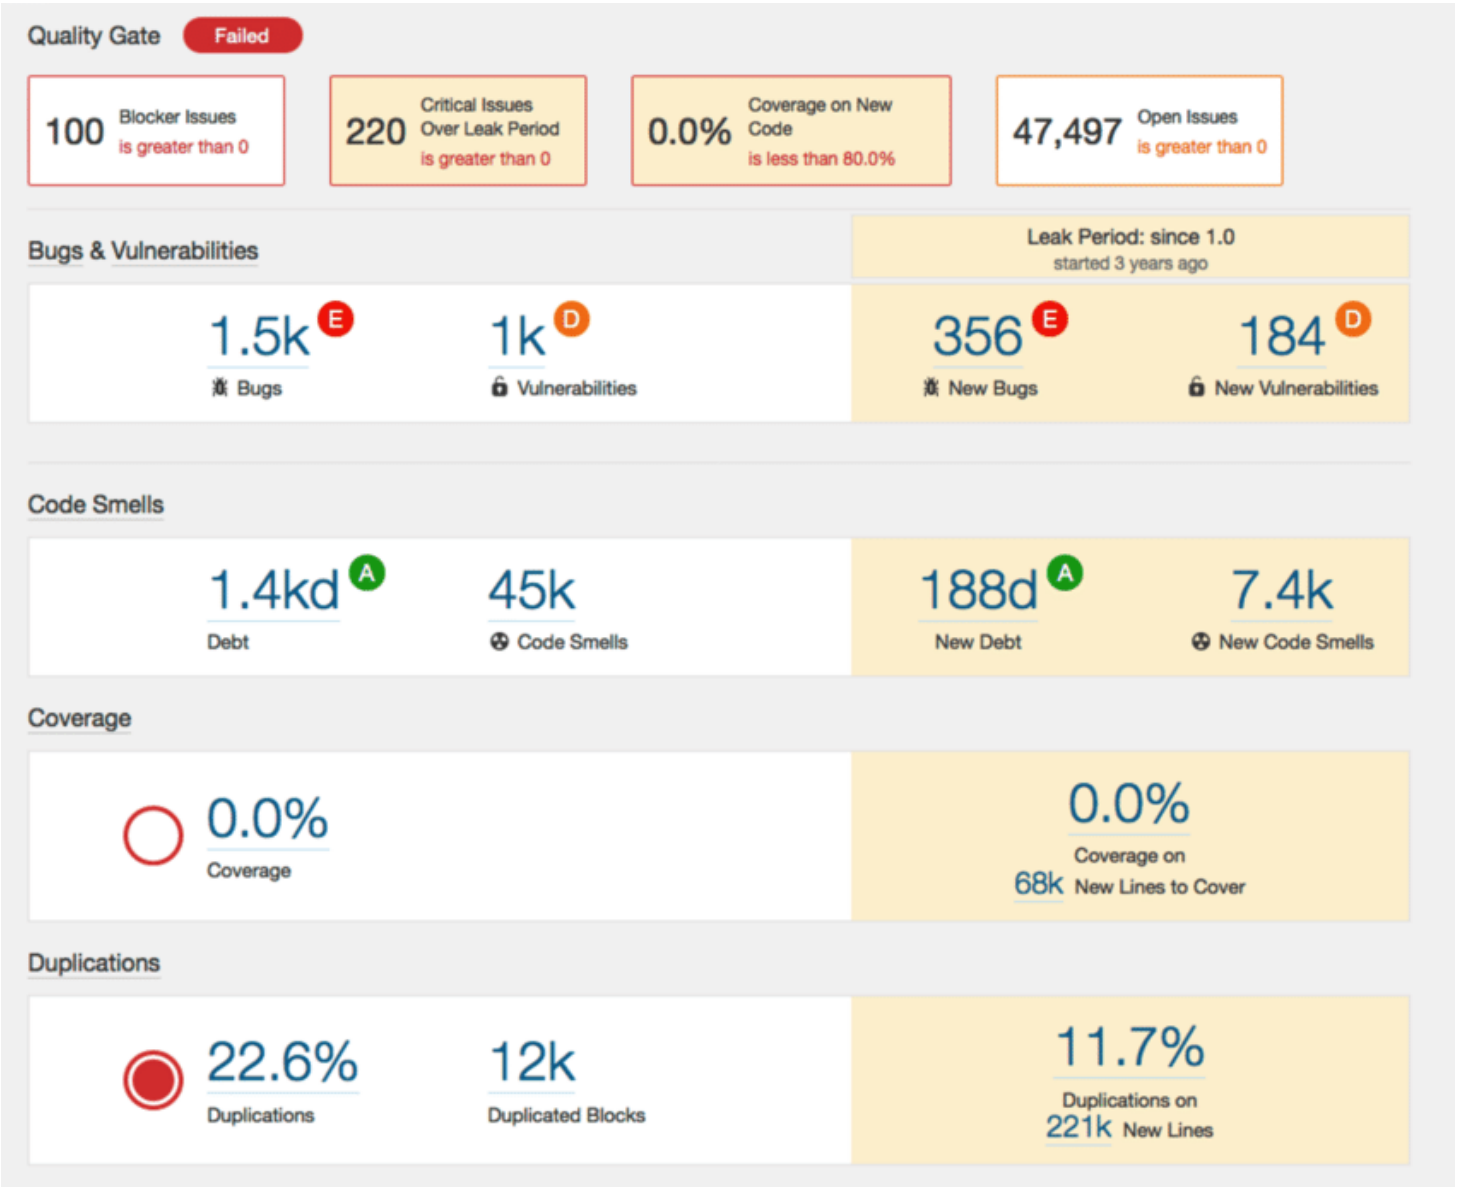
// Overview of docker-compose CLI: <http://dockr.ly/2wtQIZT>

// `docker-compose.yaml` reference: <http://dockr.ly/2iHUpeX>

// Docker Compose and Windows: <http://bit.ly/2watrqk>

WRAP UP DOCKER-COMPOSE - SONARQUBE

# SONARQUBE



# EXERCISE INSTRUCTIONS

// Setup a Sonarqube server that listens on port 9000

// Connect it to a persistent database

i.e. if you , `docker rm -f` your Sonarqube container and run a new one, no data is lost

// Use postgresql and persist it's data on the host filesystem using volumes

// Verify e.g. by creating a user via Sonarqube UI, remove the container and run a new one – is the user still present?

// Check that Sonarqube is really using your postgresql database

// **Hint:** use docker-compose

17M 43S

# SOLUTION

```
version: "2"
services:
  sonarqube:
    image: sonarqube
    ports:
      - "9000:9000"
    networks:
      - sonarnet
    environment:
      - SONARQUBE_JDBC_URL=jdbc:postgresql://db:5432/sonar
    volumes:
      - sonarqube_conf:/opt/sonarqube/conf
      - sonarqube_data:/opt/sonarqube/data
      - sonarqube_extensions:/opt/sonarqube/extensions
      - sonarqube_bundled-plugins:/opt/sonarqube/lib/bundled-plugins

  db:
    image: postgres
    networks:
      - sonarnet
    environment:
      - POSTGRES_USER=sonar
      - POSTGRES_PASSWORD=sonar
    volumes:
      - postgresql:/var/lib/postgresql
      - postgresql_data:/var/lib/postgresql/data
    networks:
      - sonarnet
    driver: bridge

volumes:
  sonarqube_conf:
  sonarqube_data:
  sonarqube_extensions:
  sonarqube_bundled-plugins:
  postgresql:
  postgresql_data:
```